

Masters Research Handbook

Lucia Rapanotti and Jon Hall

Branch: edit-stage4

Hash: 2ae634a01865baf9b439a97c67a7f9fc5ba948ca

Contents

Contents	1
0.1 Generating raw research data	6
0.1.1 Sampling: what, who (and how) to choose	7
0.1.2 Modern standards	9
0.2 Data generation methods	10
0.2.1 Observations	10
0.2.1.1 Procedural considerations	12
0.2.1.2 Other things to think about	13
0.2.1.3 Further reading	14
0.2.2 Questionnaires	14
0.2.2.1 Procedural considerations	15

0.2.2.2	Other things to think about	16
0.2.2.3	Further reading	17
0.2.3	Interviews	17
0.2.3.1	Procedural considerations	19
0.2.3.2	Other things to think about	21
0.2.3.3	Further reading	22
0.2.4	Focus groups	22
0.2.4.1	Procedural considerations	23
0.2.4.2	Other things to think about	24
0.2.4.3	Further reading	25
0.2.5	Delphi method	25
0.2.5.1	Operational considerations	26
0.2.5.2	Other things to think about	27
0.2.5.3	Further reading	27
0.2.6	Journaling	27
0.2.6.1	Procedural considerations	28
0.2.6.2	Other things to think about	28
0.2.6.3	Further reading	29
0.2.7	Fieldwork	29
0.2.7.1	Operational considerations	30
0.2.7.2	Other things to think about	31
0.2.7.3	Further reading	32
0.2.8	Documents	32
0.2.8.1	Procedural considerations	33
0.2.8.2	Other things to think about	34
0.2.8.3	Further reading	35
0.2.9	Computational thinking	35
0.2.9.1	Computational thinking tools	36
0.2.9.2	Workflow	37
0.2.9.3	Other things to think about	38
0.2.10	Mathematical thinking	38
0.2.10.1	Mathematical thinking tools	39
0.2.10.2	Workflow	40

0.2.10.3 Other things to think about	40
0.2.11 Statistical thinking	41
0.2.11.1 Workflow	42
0.2.11.2 Other things to think about	42
0.3 Managing your raw data	42
0.4 Analysis methods	43
0.4.1 Using tables to analyse data	43
0.4.2 Statistical analysis	47
0.4.2.1 Descriptive statistics	48
0.4.2.2 Inferential statistics	55
0.4.3 Qualitative analysis	60
0.4.3.1 Coding qualitative data	61
0.4.3.2 Presenting qualitative data	63
0.5 Writing up your analysis	64
0.6 Interpreting and evaluating data	65
0.7 Drafting an abstract for your project	66
0.8 Reflecting and reporting in Stage 4	67
0.9 Takeaways	68

Stage 4: Generating and analysing data

You've now reached Stage 4, which means the end of your project is now in sight. In this stage you will be in the midst of your data generation and analysis, which is possibly the most exciting, yet demanding, part of your research: this is where you get your opportunity to make that contribution to knowledge.

This stage assumes that you have worked out most of your research design details and are now in a position to begin your data generation and analysis[•].

With reference to our 5-stage framework, the activities which are in focus in Stage 4 are summarised in Table 0.1, which also provides some guidance for your interaction with your supervisor during this stage.

Activity: Understanding the effort needed in this stage

#1

Consider Table 0.1 carefully, paying particular attention to the entries in the 'Effort' column. Make a note of the activities which are most prominent in this stage and what their deliverables and learning outcomes are.

Discussion

Generating and analysing evidence will constitute by far your major effort in this stage (50% of study time): in particular, the framework assumes that you will have worked out the details of your research design in Stage 3, so you can focus on applying your data generation and analysis methods. You will also start to interpret your findings, an activity you will complete in Stage 5.

- If that's not the case, then, you should go back to Stage 3. You should also discuss your progress with your supervisor, revisiting your project timescale and risk.

Note that your data analysis and interpretation may also prompt you to generate more data, including, perhaps, reviewing more academic literature or even re-thinking or adjusting your aim and objectives better to reflect your improving understanding. Therefore, so you should expect some iteration back to activities

Table 0.1: Research activities addressed in Stage 4 (20% of project length)

Research process activities	Deliverables	Learning Outcomes: by the end of this stage you will:	Ef- fort interaction with your supervisor
Identifying the research problem	Research problem statement, refined as needed	be able to assess and improve your research problem statement	1%
Reviewing the literature	Substantial draft of your literature review, refined as needed	be able to assess and improve your current draft	1%
Setting your aim and objectives	Aim and objectives, refined as needed	be able to assess and improve your aim, objectives and related tasks	2%
Developing the research design	Research design description, refined as needed	be able to describe data generation and analysis procedures in detail	2% Suitability of methods and procedures
Generating and analysing evidence	Raw data appropriately organised and stored; data summaries and outcomes of data analysis	know the difference between various sampling approaches; be able to organise and store your raw data; be able to apply appropriate data analysis methods; be able to present your data and evidence in a concise and effective way	50% Appropriateness of data analysis and presentation
Interpreting and evaluating findings	Draft summary of findings from data/evidence generated	be able to derive findings from your data analysis and critically assess them in relation to research aim and objectives	15% Critical and logical thinking
Reflecting and reporting	Stage 4 report; draft abstract for your project	know the purpose and content of an abstract; be able to assess your research progress and write up a substantial report, including an abstract for your project	25% Any further improvements required, particularly in relation to critical thinking and academic writing
Planning work and managing risk	Updated risk and work plan	be able to assess risk and draw a work plan	5% Any major adjustment required to address deficiencies or manage risk

you have carried out in previous stages, and revision of things you have written.

0.1 Generating raw research data

Your *raw data* represent any data you generate as part of your research. Which data you generate, and how, is determined by the choices you have made in your research design, informed by your research aim and objectives. In this chapter, we look at key methods for generating data, by providing:

- design tips and techniques for each method;
- a basic workflow to get your data generation going;
- weaknesses exposed through the data generating technique;
- further sources to consult for more detail.

When we say “data generation”, we don’t necessarily mean that you will create a new-to-the-world “data”[•]. We simply mean new-to-your-masters-project data, which covers a multitude of sins, including:

- the creation of a brand new data set, which did not exist before your research project. This may be the result of data collected through new observations or measurements, a new survey, questionnaire or focus group, through selecting passages from documents, etc.
- the extension of a previously collected data set with new elements, derived from those that already exist, for instance adding the mean value of a collection of numerical data, or grouping together specific distinct data into new categories that you have created;
- the collection of previous data for reinterpretation, for instance if you are rerunning a previous experiment in order to confirm its results, or doing a meta-analysis of the literature in a particular area.

[•] Although this might indeed be the outcome of your data generation.

Our use of the term data generation includes all of these. Of course, in making a contribution to knowledge, you’re going to have to do it regardless of the data generation method and so the data you generate as part of your research must allow you to conclude something new.

Related to data generation is the concept of “data source”, which is the location from which your data originates. If you are re-using existing data sets, this may well be an archive or a digital repository. For new data sets you generate, this may well be the experimental or real-world setting of your own observations and measurements, or a population of interest from which you will derive a “sample” for further analysis. The latter underpins many data generation methods, so that we will consider sampling in some detail next.

0.1.1 Sampling: what, who (and how) to choose

Sampling is the process of selecting a subset[•] for further analysis from an infeasibly large population of interest.

Sampling assumes that there is a “sampling frame” as data source: a sub-set of the collection of the population of interest from which your sample is taken. In some sense, you have already experienced sampling as part of your literature review[•]. Unless you had infinite amounts of time – which you didn’t – and infinite patience – which you might have – you could never be 100% certain that your literature search collected *all* relevant papers: the search space is infeasibly large (and not indexed particularly well). But you were systematic and achieved a practically good[•] coverage because of that.

More generally, sampling is used when we wish to study a population of interest, but this is infeasibly large or inaccessible for us to be able to study every single member of it. Instead, we choose a sample which is somewhat representative of the population characteristics, hoping that by studying the sample we can establish some properties or patterns of interest which can be assumed true of the population as a whole.

Broadly speaking, sampling can either be random or non-random.

In *random sampling*[•] some unbiased way of choosing the subset members from the population must be established upfront to inform sample collection, which is usually completed prior to any analysis. This is used particularly in quantitative research, and when generalisation of the results to the population is of primary importance, something enabled by the lack of bias in the sample selection process.

Instead, in *non-random sampling*[•], the sample choice is based on the researcher’s judgement and discretion, so that some element of bias may exist. Members can be added to the sample as the research progresses, interleaving data collection and analysis, until no more collection is possible or *saturation* is reached, that is collecting more data would not bring more relevant information. This kind of sampling is used particularly in qualitative research, where depth and richness of results are more important than the ability to generalise.

Random sampling techniques include:

simple random sampling where each member of the population has exactly the same chance to be se-

- We could have said “sample” but that would have been circular.
- You might remember the relatively complex procedures for recording search terms, discovered papers, their relationships, and your growing collection of notes on them.
- By *practically good*, we mean you found the most of the most important papers, some other papers, and didn’t have to read *every single paper*. I.e., you found a *representative sample*.
- Random sampling is also called *probability sampling*.
- Non-random sampling is also called *non-probability sampling*.

lected. It has the advantage that it is easy to implement, and given the complete randomness of the sample, generalisation is fairly reliable. However, it can be time consuming if the population is very large, and may not lead to a representative sample if the population has large sub-groups, which may be over-represented in the sample, with minority groups being under-represented.

stratified random sampling where sub-groups of the population are identified based on common characteristics, the *strata*, and sampling is random across those strata. The strata are not mutually exclusive: for instance, the population may have sub-groups defined by gender, ethnicity and level of education, which may overlap. This approach overcomes the over/under representation problem of simple random sampling; however, deciding on the strata may be difficult and will also complicate data analysis.

cluster sampling where the population is divided up in naturally occurring separate clusters, and the sample is obtained by randomly selecting some clusters and then randomly selecting members of those clusters. It is more cost-efficient than the other two approaches, but can introduce bias if the selected clusters are not representative of the whole population, so that the over/under representation problem remains.

Non-random sampling techniques include:

purposive sampling in which participants are selected by the researcher based on particular characteristics, knowledge, or expertise they have. It is often used for small, rare or unique populations, and is particularly suited to studies which intend to be deep and narrow, and for which generalisation to the population is not the main concern. As the sample choice is made by the researcher, it is prone to bias. However, it also allows the researcher to involve participants who can provide insights into such rare or unique groups.

convenience sampling where participants are selected based on their availability or accessibility. This is quick and easy, but unlikely to produce a representative sample, so, once again, bias is an issue.

snowball sampling which relies on referral from previous participants to recruit new ones. This is an effective approach when a population is difficult to access or when the topic is sensitive or tabu. This too is unlikely to generate a representative sample, and is prone to bias. However, it is a way to gain access to members of a population which may be otherwise inaccessible.

In summary, when choosing a sample, you need to consider various factors, including the aim of your study, the kind of methods you are applying, and the level of access you may have. Trade-offs are likely

involved and you may not be able to obtain an ideal sample. Nevertheless, your sample will still be useful to your research, as long as you clearly explain and justify how it was obtained and what its limitations are.

Activity: Deep reading on sampling

#2

Check back to your choice of research strategy. If you've chosen one which may require sampling, then you should go deeper into this topic to ensure you select the right kind of sampling for your study. We recommend you start by reading the following: <empty citation>

Guidance

The suggested reading is only a starting point. You should go deeper into the specific kind of sampling you are most likely to apply.

Activity: Your chosen sampling approach

#3

Assuming your study requires you to perform some sampling, write down the sampling approach you are going to apply, with its justification in terms of your aim and objectives, and any trade-offs due to the practicality of accessing the sample. Record any possible weakness or limitation of your chosen approach, and how you will address them in your project.

Guidance

You can skip this activity if sampling is not required by your choice of research strategy.

0.1.2 Modern standards

Modern standards of research often require that your data be made available to other researchers so that your research can be verified or even rerun. In fact, it is increasingly the case that data sets are published and shared by entire research communities, often used as testbeds or benchmarks for new knowledge contributions. For instance, in medical applications of Machine Learning, in which new knowledge can be the fractional improvement of the performance of an AI algorithm to, say, diagnose a medical condition from images, not being able to share the data set you have used in your research can negate your knowledge contribution. Therefore, you should consider whether your data (or a sample of it) should appear as an appendix to your

dissertation, or even whether it should made available in its entirety to the wider research community, and how.

Modern standards of research also require you to comply with regulations on data privacy and protection[•], so that as part of your data generation process you should also consider the need to anonymise data[•] without losing their research value for your project or the need to release commercial in confidence or otherwise sensitive data.

0.2 Data generation methods

In this section, we step through the most used data generation methods – those that were mentioned within the research strategies in Stage ??, i.e., Interviews, Journalling, Observations, Questionnaires, Documents, Focus groups, Field work, Computational thinking, Mathematical thinking, and Statistical thinking.

Most of these methods concern *empirical data*, that is data that is gathered through our five senses or from experience, and then used as the benchmark against which theories and advances in knowledge are made.

0.2.1 Observations

Observations constitute one of the main ways in which empirical data are generated[•]. In fact, according to **marvasti2014analysing**, “Observation is the foundation of science.”

This method requires the researcher to make their observations[•] of a phenomenon of interest. What you will observe are core characteristics of the phenomena that you have identified as part of defining your research problem.

Activity: Do I need to know about observations?

#4

Check back to your chosen research strategy from Stage 3. Does it involve data generation using observations? If so, read through this section and complete the activities.

Observations can be made directly, through our naked senses, or through instruments which enhance our sensory capabilities, such as a telescope, a microscope or the “a myriad of other ingenious inventions designed to make the invisible visible, the evanescent permanent, and the abstract concrete” **daston2011introduction**. Quantitative observations, say, the size or weight of an object, are usually referred to as ‘measurements’.

- We discussed GDPR in Section ??.
- There may be a route to your degree in which your thesis is not published. Should you have any concerns about the release of data, please do find time to discuss this with your supervisor, balancing the needs for total anonymity with those that can be achieved through data anonymisation.

cross-check with stage 3 if this is the full list

- Whole books have been written on observations as a research method; we can give but a shallow introduction. Further sources you can use are included at the end of this section.
- Hence the name...

What to observe Observations are versatile tools for almost any research domain, and your own domain will determine what sort of observations you will make. Observations range across natural phenomena – such as the different proportions of plant species that populate a train station wilderness garden – through the artificial phenomena – the way that buses drop off and pick up their passengers at a train station – to observations of social phenomena – the different ways in which a train station is used by commuters in the morning and the evening – as well as more complex combinations of each. Each will use different observations techniques and tools, and each with different constraints.

Observation types Observations can be *naturalistic*, when phenomena are observed as they happen in their natural setting – for instance, observing the behaviour of animal species in their habitat, or *structured*, when phenomena are observed in a somewhat artificial environment, such as during an experiment – for instance, giving people a specific task to perform and observing how they carry it out. In this case, often the aim is to collect quantitative data, say the speed at which the people can complete the task.

More on social observations As an observer of people, you can act as either a *participant* or a *non-participant* observer. The former is a researcher that interacts as a member of a community under observation, becoming an active participant in the group or situation under study. In effect, as a participant observer, you would be “living” alongside those you observe – you might be a commuter that uses the train station and so share the experience of the other commuters you observe. Instead, in non-participant observation, you would remain separate from the group or situation being observed – you may observe commuters using the train station, but not actively become one of them.

Depending on whether people are or not aware of being observed, observations can be *covert* or *overt*[•]. Covert observations have the advantage that people’s behaviour is not affected by their awareness of being observed, but, of course, they raise some important ethical and legal issues, in relation to informed consent, and privacy and anonymity. If the judgement is that the phenomena observed do require privacy – perhaps you wish to observe commuters’ use of restrooms in the train station or customers in a betting shop – then you must explicitly ask for permission – or change your research problem! Otherwise, in a public space, there may be no overriding expectation of privacy and observations can be done without explicit consent. Your university is likely to have strict regulations on the matter, or even prevent you from conducting covert observations as part of your research.

[•] The terms disguised and undisguised are also used in the literature.

0.2.1.1 Procedural considerations

In order to apply this method, some preparation is needed for you to decide what you will observe as how. Specifically:

The phenomena You will need to decide which phenomena to observe, whether natural, artificial, or social: this choice will depend on your research problem, and aim and objectives.

The kind of observations Depending on the phenomena, you will need to decide whether you will perform naturalistic or structured observations. In addition, for social phenomena, you will need to decide which mode, whether participant, non participant, overt or covert. For the latter, you must also identify the steps you will take to ensure compliance with ethical and legal guidelines.

The time and place For all kinds of observation, you must determine the time and place at which those observations will be made. For participant observations, this choice will be determined by your participation in the group or activity being observed, which may or may not be under your control. For instance, if you are a participant observer of train station usage, then you can determine when to use the station and make your observations. However, if you are a participant observer in a change project within your organisation, then the timeline of the project will determine when your observations can take place. In all cases, you should draw a schedule which establishes the timing and frequency of your observations, and which provides an efficient way for you to conduct your observations. For instance, it may be that you make exaggerated use of the train station[•] to condense many months of participant observation into weeks or days – after all, your research project is time bounded – perhaps visiting ten times per day rather than just two.

The use of instruments You may be able to make your observations using only your senses. However, many phenomena do not permit observation through the senses unassisted – for instance, the search for exoplanets[•] requires complex and delicate instruments which will be located on mountain tops. In this case, the availability of the equipment you will use will determine when and how you make your assisted observations — you will also need to gain access to the equipment, and this will constrain your schedule and may alter your research plans[•].

How to record your observations It is important to record what you directly observe, separate by any added interpretation[•] to avoid possible bias affecting recorded observations. To this end, observations are

- You could eat there, for instance, or use any shops that are colocated.
- For instance, **jones2008exoplanets**. But don't let this exciting mission to explore strange new worlds; to seek out new life and new civilizations; to boldly go where no one has gone before distract you. Too much.
- Remember, plans never survive first contact with reality, so also plan to have a backup plan that you can use should your first plan to make your observations fail!
- That should happen later on, as part of your data analysis.

typically recorded in notebooks with a double entry, which separates pure observations from possible value judgements made by the observer – for instance, observing someone “waiting impatiently for the train door to open” is ascribing feelings to the person observed which, by their nature, are hidden from the observer, but may be inferred from the observed’s body language. In this way, another researcher reading your notes can clearly differentiate direct observations from such inferences. In some cases, you may be able to use audio and video recording, say using your smart phone, to capture your observations for follow-up analysis. In this case, ethical issues in relation to privacy and informed consent also apply.

0.2.1.2 Other things to think about

Observations are by no means an easy way of generating research data, and there are many issues that can arise, including:

Hawthorne effect Overt observations can lead to the so-called Hawthorne effect[•], which consists of people changing their behaviour due to their awareness of being observed. This can be mitigated by building a rapport, including spending more time with the people being observed and observing them for longer periods of time, and, in case of structured observations, by ensuring that the tasks participants are asked to do come natural to them.

Observer bias All observations can be influenced by the observer’s own bias, whether implicit or explicit, or the result of overfamiliarity with the phenomena of interest. To guard against it, triangulation should apply, including using different data sources and collection methods, or having multiple observers all following a standardised procedure. The use of double-entry notebooks, as described above, can also help, as they separate pure observations from interpretation and inferences made by the observer, with the latter being the subject of scrutiny for potential bias.

Volume of observations Observations can lead to vast amount of data to analyse. Different analysts, or different stages of analysis as your research progresses, may focus on different aspects of the same data. This can be a good thing should you analysis deepen due to understanding more about your observations, but may also lead to analysis drift or even “paralysis through analysis” in which no progress is made due to too much depth. To avoid this, keep a clear eye on the prize: your research goal, and set regular times at which you can reflect on progress.

- This term was coined in the 1950s in relation to a productivity study carried out at the Hawthorne Works, and electric plant new Chicago.

0.2.1.3 Further reading

Activity: Deep dive into observations

#5

To find out more about observations, take a look at these resources:
daston2011introduction;
marvasti2014analysing;
simpson2003using;
driscoll2011introduction;
angrosino2003observations;
sapsford1996data

0.2.2 Questionnaires

Activity: Do I need to know about questionnaires?

#6

Check back to your chosen research strategy from Stage 3. Does it involve data generation using questionnaires? If so, read through this section and complete the activities.

Questionnaires[•] are versatile tools for generating data from participants by asking questions[•]. They allow a researcher to collect participants' answers about their attitudes, preferences, opinions, behaviours, etc. You might use a questionnaire as a way of collecting statistically significant responses from a population sample, but there are other uses as well, for instance as the basis of interviews[•].

Activity: Do I need to know about questionnaires?

#7

Check back to your chosen research strategy from Stage 3. Does it involve data generation using questionnaires? If so, read through this section and complete the activities.

If you do use a questionnaire, its thoughtful design is of critical importance. Otherwise, you might be asking your (willing) respondents to spend a considerable amount of their valuable time answering questions

- Questionnaires are just one in a rich collection of *survey tools*, others of which are described below.
- There's a hint in the name – *questionnaire* – although why two “n”s; does no millionaire, billionaire, or debonaire use them?
- Which we cover in the next section.

the content of which are not helpful for your research. As they might not be so willing to help a second time, getting the questions right[•] the first time is important.

Administering questionnaires are nowhere near as difficult as they used to be as the number of online resources for doing so increases. And, probably because of this, there are plenty of resources in the literature and online to help you design your questions. Their descriptions can be a little technical, however, so the following glossary might help you engage with them better.

Essential questions the smallest possible set of questions you absolutely need to ask to address your research aim and objectives. While using several questions will give you richer data sets, long questionnaires tend to put people off, so that fewer people may be willing to participate.

Profiling questions questions that ensure your respondents match specific characteristics you are interested in: say, you are studying the usability of a new product, then you will need to know the extent your respondents have engaged with that product. This is particularly the case if you are running a large survey and don't know who is going to respond.

Demographic questions often used so that you can then compare answers across different sub-groups, say, based on gender, age or ethnicity, etc.

Response options questions are broadly divided into *closed* and *open-ended*. Close questions restrict the possible responses to a set of given choices, while open questions allow respondents to use their own words freely to answer the question.

0.2.2.1 Procedural considerations

Using tools While you can design your questionnaires from scratch using your word processor, there are plenty of specialised digital tools, many of which are free, that can make it a lot easier[•]. They usually come with: templates and pre-defined question types that you can customise for your study; statistical analysis and data visualisation features that you can apply to the data you have collected; export functions that allow you to save the data to a spreadsheet for further analysis. Overall, if you need to develop questionnaires for your research, they can really help you speed up the process, so that it's well-worth the investment of time in climbing their learning curve.

- Often called *questionnaire design*, although this conjures up glossy format and whizzy web-pages which is of secondary importance. Unless, your questionnaire is about the design of questionnaires, of course.

- Examples include: [add list here](#) and [URLs](#).

Drafting, testing and piloting Unless you have many years of experience in questionnaire design, your first questionnaire draft will be far from suitable. Indeed, releasing your first draft without further thought may lead to you not only generating no useful data from it, but also putting off your audience sufficiently that they are not willing even to look at your second version. So, once you have a first draft of your questionnaire, you should test it and refine it.

Early testing can be done by asking a friend, a family member[•], or a colleague to work through the questions, provide their answers and any other feedback they might have. This will give you early indications of problems with your questionnaire[•], which you can spot, for instance, if the respondents are confused – which may point to a lack of clarity in the questions – or hesitant – which may point to a poor choice of response options or to inappropriate scales – or disengage before completion – which may point to too many questions being asked. Be sure to loop back to those that have helped you to check that you have addressed their comments.

Later in the process of designing it, however, you should take expert advice including, of course, that of your supervisor, to get to the final agreed form. In addition, you could pilot your questionnaire on a small number of respondents first, then revise it as necessary before using it more widely.

0.2.2.2 Other things to think about

In designing your questionnaire you should pay particular attention to the following:

Plain language Your questions should be clear and plain, and you should avoid jargon and idioms, to ensure your participants understand what you are asking, particularly if not native speakers.

Unbiased language Your questions should also be objectives, that is you should avoid any judgemental term or tone which may reveal your own opinions or beliefs, or influence participants to answer in a particular way. You should also avoid questions which make assumptions about your respondents' habits or behaviours: for instance, asking participants what they eat for breakfast, assumes they all take breakfast, which may not be the case.

Double-barrelled questions Also termed “compound”, these are questions that ask more than one thing, while only allowing one answer. These should be avoided as it would be difficult, if not impossible, to establish in your analysis which part of the question each participant has answered. Instead, you should split the question into separate questions each addressing a specific thing.

- Probably, but not always a friend:)
- Although it's sometimes difficult, you'll make more progress and quicker if you think of the questionnaire as imperfect, rather than you. You can then apply comments – even if they are negative – to the questionnaire rather than having a personal emotional reaction to them. For each comment, make sure you understand how it can be addressed in your questionnaire. This last tip also means that you can welcome (but ignore) comments that can't be addressed.

Closed and open questions For your closed questions you should ensure that the possible answers provided cover all possible options[•] and do not overlap, that is they are mutually exclusive. Instead, for your open questions you must ensure they are sufficiently constrained so that your participant's answers don't end up being too vague or off topic, hence not providing much value to your research.

Scales If your questions require participants to estimate or measure something, you need to worry about both validity and reliability when setting up the scales for possible answers. In this context, validity means that the chosen scale should allow respondents to measure something accurately; while reliability means that, under the same conditions, respondents will be able to come up consistently with the same (or very close) measurements.

Question grouping, ordering and flow You should group related questions[•] together, and establish a logical flow in sequencing groups of questions, so that topics follow naturally from one another. Question order in each group also matters: as a rule of thumb, simpler questions should precede more complex ones.

- Or, at least those you're interested in.

- For instance, those intended to establish a demographic of respondents.

0.2.2.3 Further reading

Activity: Deep dive into questionnaires

#8

To find out more about questionnaires, take a look at these resources:
oates2008researching;
hays2003case;
burns2009action; **mcclure2002common;**
najafi2016observation;
robertson2002automated;
kielmann2012introduction

0.2.3 Interviews

Interviews are a method for generating data from participants by asking questions and recording detailed answers. They are a form of conversation between the researcher and one or more interviewees, designed

by the researcher to gain insights and opinions on a specific topic. The researcher guides and controls the conversation and asks the questions.

Activity: Do I need to know about interviews?

#9

Check back to your chosen research strategy from Stage 3. Does it involve data generation using interviews? If so, read through this section and complete the activities.

The personal approach that is a characteristic of interviews means that they are a great way of accessing a (group of) individuals' feelings, thoughts, ideas, and/or experiences, data that can be difficult to generate in other ways. They can help you obtain detailed information on a specific issue or topic, asking open-ended questions which may be tackled or interpreted differently by different interviewees. They are also an effective way to investigate sensitive issues or privileged information that interviewees may not be willing to commit to writing.

Interviews can also provide direction for new research by giving expert indications of where problems lie in a particular domain. As such, interviews can often be used as a way into a discipline, filling in useful background through personal experiences, and having access to otherwise difficult to access information.

There are three main kinds of research interviews:

The structured interview serves as a repeatable framework by which each participant is asked the same questions in the same way. There is no scope for deviation from the structure, so that auxiliary questions and follow-ups are not used.

Structured interviews are the closest to being time bounded and predictable; if you only have 8 hours to conduct 24 interviews for instance, a structured interview would be the best way to achieve this.

Your skills as an interviewer will be tested by structured interviews: it is often difficult not to stray outside of the structure when an interesting answer is given, and you may have to cut a participant short if their answers overrun or diverge from the structure[•].

The semi-structured interview serves to identify areas of interest to the researcher, with interesting responses being welcomed and followed up if appropriate. Interviewing a domain expert on your chosen topic would be well-served by a semi-structured interview as their expert knowledge could be probed with follow up questions.

- In our experience, this is a perennial problem, so don't underestimate the difficulties you will face as an interviewer.

The semi-structured interview does not naturally time-bound the interaction, and so – if you don’t have unbounded amounts of time – you will have to balance the breadth of questions with the depth of responses.

The unstructured interview has no structure to constrain the route through the data that the interviewee wishes to take, although the interview may begin with the same question each time. As the direction may wholly be decided by the participant, you challenge may be retaining forward motion and focus during the interview.

0.2.3.1 Procedural considerations

Main factors to consider in interviews is to ensure that you have chosen appropriately those who will take part, the kind of interview and the questions you will ask and their structure, the tools that you need to take full value from the data generation, and how you will followup. Interviews do not need to be co-located – a video or audio link is all that is needed – but they do need to be in real-time.

Interview type In being standardised, structured interviews make it easier to compare your interviewees’ answers objectively. However, interviewees are limited to the set questions, so there is no scope for digging deeper into their answers. If you need deep insights, you should use the less structured interview forms which allow you to probe your participants responses.

If you have a good level of domain knowledge, you can use the less structured interview models, as you will be able to follow your participants’ answers more easily and direct their comments towards your research interests[•]. If you don’t already have a good level of domain knowledge, then you will be putting more effort into the design of the interview, so that you can simply capture your participants’ responses to your – well-designed – questions.

- Of course, you can always use structured interviews even if you do have domain knowledge – there’s nothing to stop you.

Whom to interview You will need to choose your interviewees carefully. In the perfect case, you should interview until the responses you are receiving are “guessable”[•]. Practically, you will have limited time and resources, and limited access to interviewees, so that sampling may be required: in such case, you should follow the advice on sampling in Section 0.2.3.

- I.e., you no longer get novel answers to your questions, indicating that the topic has been covered.

Ethical and legal matters Your university will have strict guidelines on how to approach and work with human participants, which you should investigate before you contact your potential interviewees. At

a minimum, those guidelines will cover informed consent, handling personal data, and health and safety, but they may also prevent you from interviewing certain groups of people, for instance minors or vulnerable adults. You should go back to the advice in Section ?? to refresh your understanding of ethical and legal issues in research.

Testing and reviewing You can apply some of the advice in Section 0.2.2 in relation to designing your interview questions. Once you've drafted your interview questions, you should do a dummy run of your interview with a willing friend, family member or colleague, and use their feedback to improve your questions*. In particular, you should consider:

- whether you were able to put the participant at ease during the questionnaire. If not their nervousness might influence their ability to contribute, and you should consider what to do differently
- which questions worked well and led to useful responses, and which were confusing or led to unhelpful answers. For the latter, you should consider how to reword them – perhaps with the help of your participant*, by having alternative versions of the questions, or by removing or replacing the questions
- (if time constrained) which questions overran, and whether you can rephrase them to be less “open”
- (if structured) whether you were able to keep to your “script”. You should reflect on how will you resist the temptation to probe more deeply, or whether you should consider moving to a semi-structured or unstructured form
- whether the questions were in a logical order, ideally they grouped by topic. If not, consider re-arranging them to build responses in the most productive way
- whether you have sufficient questions to elicit the data you need, or there other questions you should ask

Once you are satisfied with your questions, you should run them past your supervisor – they will be sure to have comments.

Recording answers You need to decide how you will capture your interviewees's answers. If you are planning to use audio or video recordings, you will need to ensure your participants are aware and give their explicit consent. If not, you will need to take notes manually. In this case, you should also test your note taking, to ensure you are able to capture everything of interest*.

* Repeat this with as many willing participants as you can until you're happy with the interview format or until you run out of willing participants, or time! Make sure that you do not dip into your target population for these preliminaries.

* “I would have asked it this way” ...

* Longhand notes can be taken at 35 words per minute; spoken text is often as fast as 120 words per minute.

Where to hold the interview With the advances in video conferencing technology, interviews can be conducted effectively online, with the added bonus that they can be easily recorded, often with transcripts automatically generated. However, interviews in a physical space where you are colocated with your interviewee, remain common. For these, you will need to ensure that you have an appropriately comfortable venue for your interview, including access to comfort facilities. Public spaces – where you could share a coffee, for instance – may create a more immediate feeling of intimacy, and so deeper responses, but they may not be suitable if discussing sensitive issues or if background noise might interfere with your record keeping. Therefore, make sure to check the venue out at the appropriate time of day to ensure it is appropriate for the interview and a recording device can handle any difficulties.

Opening and closing interviews As part of giving their informed consent, interviewees should be fully aware of what you are trying to achieve in your research and what the purpose of the interview is within it. They will also be interested in how you will use their answers, and should be reassured as to any use of confidential information or personal data. It is therefore good practice to provide this information at the start of your interview, or even as part of inviting them to participate: perhaps share a sheet which includes this information and describes the research you are doing. At the end of the interview you should also thank them and explain what will happen next, including how they can get in touch if they have any further concerns and follow up questions.

0.2.3.2 Other things to think about

Making a checklist There is lots to think about when preparing and conducting interviews. Have you made all necessary arrangements to conduct the interviews? Did you obtain all necessary permissions, including informed consent? Do you know how you will record the answers? What would happen if your audio recording device went wrong? Would you have a backup or the interview? What if you forgot to turn it on?• Write a checklist of instructions for yourself to follow before and after each interview, so that you can be sure not to miss anything important.

- Oh so easy to do...

Being a good interviewer Try, to the extent possible, given the format, to allow your participant to govern the speed and direction of the interview. Allow them to talk in complete sentences without interruption, or have a good reason to interrupt. If you need to interrupt, apologise for doing so and tell them the reason why you have done so*. Be polite and encouraging, as your participant might be nervous.

- “I’m sorry to have to interrupt, but we only have 5 minutes left and ...”.

0.2.3.3 Further reading

Activity: Deep dive into interviews

#10

To find out more about interviews, take a look at these resources:
[oates2008researching](#);
[johannesson2014research](#);
[secor2010social](#);
[hays2003case](#);
[mcclure2002common](#);
[peoples2020write](#);
[jorgensen2001grounded](#);
[hycner1985some](#); [englander2012interview](#); [ramsook2018methodological](#);
[robertson2002automated](#);
[kielmann2012introduction](#)

0.2.4 Focus groups

Focus groups engage participants in interactive discussions to develop an understanding of complex phenomena and generate new hypotheses for further research or practice. They are effective at surfacing a full range of perspectives held by the participants and, through their interaction, expand on their individual contributions. They are particularly useful to uncover data and ideas that may not come up in one-on-one interviews, and are generally a more efficient way of collecting data than multiple interviews.

Activity: Do I need to know about focus groups?

#11

Check back to your chosen research strategy from Stage 3. Does it involve data generation using focus groups? If so, read through this section and complete the activities.

Focus groups include participants who share some common characteristics or interest. They are moderated, often by the researcher, so that they combine elements of interviews and observations alongside the group discussion. The moderator plays a crucial role in facilitating group processes, maintaining focus, and

controlling participant interactions. Depending on the research aim, it might be necessary to run a series of focus groups so that trends across different groups can be identified.

There are different flavours of focus group. For instance, there may be two moderators with separate roles, say one looking after the procedures, the other focusing on the discussion, or both contributing to the discussion, but taking opposite sides or playing devil's advocate[•]. You could have a two-way focus group, in which there are actually two moderated groups each listening to each other's discussion, with a view to stimulate richer insights through rebuttal or further elaboration of ideas. You may also reduced the number of participants to create a more intimate 'mini' focus group.

To optimise data collection from focus groups, careful attention must be paid to the composition, size and number of groups, selection and training of the moderator(s), and development of the questions used to guide the group discussion.

0.2.4.1 Procedural considerations

Group type Depending on your research aim, you should decide which kind of focus groups you will need, including whether more than one moderator is required, or if a series of focus groups would be desirable.

Group size The number of participants in a focus group is usually from between 8 to 12 participants[•], although other sizes work too, the smallest useful size being 4 to 8 for 'mini' focus groups. Anticipating subject loss, you should over-recruit participants by approximately 25%.

Participants The purpose of a focus group is to obtain data regarding ideas, attitudes, understanding, and perceptions on a specific topic, and choosing participants that can contribute to this purpose is an important part of identifying the right participants. Participants should therefore be selected based on their experience and interest in the topic, rather than through random selection. Although the potential range of participants might be limited by context – you may need them to be selected from a small organisational group – the choice should be made so that they come from as diverse range of backgrounds, views, and experiences as possible. Participants should not, however, be chosen to be individuals suggested by fellow group members.

Moderation A skilled moderator is needed to guide the discussion. Moderators' key qualities include empathy, positive regard[•], being able to use of pauses and probes effectively in the group discussion, and

- The term 'duelling moderators' is used in this case.

- Is a 12-person jury simply a focus group?

- This denotes a general affirming caring, and supportive attitude.

exercising control in an unobtrusive manner. If you do not have access to a skilled moderator, then accessing moderator training for yourself might be desirable[•].

Location Focus groups can be run online or participants can be physically co-located. For the latter, placing participants within an uncomfortable environment is likely to lead to negative outcomes. Given that a focus group might last for an extended period, appropriate timing should also be considered, with access to comfort facilities, etc., and an explicit timetable which includes breaks.

Choice of questions Focus groups are sometime referred to as group interviews, in the sense that the moderator seeds and controls the discussion by asking questions. It is important therefore that you consider which questions to ask, including opening questions to get the discussion going, or questions to probe further and to ensure all participants get involved. Open-ended questions are the norm in focus groups as the intention is to elicit insights, attitudes, opinions and perceptions.

Discussion etiquette You will need to establish an etiquette for the group discussion, including expected participants' behaviour, for instance in addressing each other, taking turns when speaking, whether mobile devices should be switched off, etc.

Recording the discussion Usually video or audio recordings are used to record the group discussion, so you will need explicit consent from the participants. Note taking is possible but only if there are separate moderators and recorders.

0.2.4.2 Other things to think about

Groupthink The point of a focus group is to elicit diverse views from participants, so it is important to be wary of *groupthink*, a tendency to conform to majority opinion to maintain unanimity and avoid confrontations, and which may inhibit discussion and the expression of diverging views. As a moderator you can mitigate against groupthink by asking probing questions, ensuring that a plurality of views are expressed, or playing devil's advocate in relation to prevailing ideas.

Social desirability bias This is the tendency of participants to express opinions which they think are more likeable or acceptable by the group, even if they are not honest accounts of their views or experiences.

- There are videos purporting to guide moderators on youtube.

The moderator can mitigate against this bias by framing a question in an hypothetical or indirect manner, to distance it from the participant's personal experience, the latter being something they may be reluctant to share. Establishing an atmosphere of trust, anonymity and confidentiality can also help participants being more open and honest.

Group dynamics Group culture and power relations, and participants' personality may also introduce bias and affect the end result. For instance, shy participants or introverts may feel overpowered and intimidated by assertive participants, whose views may then become prevalent. The moderator has the task to ensure all voices are heard, possibly by calling out shy participants individually, or time-limiting contributions to prevent the most talkative participants from taking over. Larger groups may be more difficult to manage and control, so that group size should be chosen wisely.

0.2.4.3 Further reading

Activity: Deep dive into focus groups

#12

To find out more about focus groups, take a look at these resources:

powell1996focussmithson2000usingplummer2008focus

<https://www.eiu.edu/ihec/Krueger-FocusGroupInterviews.pdf>

0.2.5 Delphi method

With the Delphi methods[•], a group of experts are consulted individually by the researcher with a view of obtaining a consensus on a particular issue, problem or topic.

Activity: Do I need to know about the Delphi method?

#13

Check back to your chosen research strategy from Stage 3. Does it involve data generation using the Delphi method? If so, read through this section and complete the activities.

The method involves an iterative process of collecting, synthesising and circulating anonymous judgements among those experts to arrive eventually at a consensual view. At each iteration the experts can revise

- Also called Delphi technique in the literature. Its name is a reference to the ancient Greek temple that hosted the Oracle of Delphi, famous for her prophecies.

their opinion in light of what has emerged in the previous iteration. Anonymity is used to ensure that no individual expert exercises undue influence on the other experts, hence also mitigating against groupthink, and all participants feel free to express their opinions without any fear of judgement or criticism, hence also mitigating against social desirability bias.

This method is based on the assumption that a group of experts are more likely to arrive at an informed and valid position than an individual, due to the diversity of knowledge and experience. The judgements and consensus gathered constitute the generated data that are subsequently analysed by the researcher.

The Delphi technique is suited to situations where it is important to access collective expertise, particularly to inform decision making, policy creation, risk management or forecasting.

0.2.5.1 Operational considerations

Selecting experts Participants are selected based on their knowledge and experience in relation to the issue, topic or problem under study, so this is purposive rather than random sampling. Diversity of experts is important to ensure breadth of expertise, hence to generate valid outcomes. Between 10 and 50 experts are usually selected to participate in a Delphi study, although both smaller and bigger numbers have been used in studies reported in the literature. The more participants, the more resource-intensive the process of collecting, analysing and combining feedback is going to be.

Process Maintaining anonymity is essential throughout the process. Each expert is initially consulted separately by the researcher, who then anonymises and aggregates the group responses and circulate them to the same group of experts to seed the next round of consultation. Providing feedback at each round is the essential mechanism to foster convergence of opinions, as such feedback is used by each expert to review and refine their own opinions or judgements. In theory, this process is repeated over multiple rounds till a consensus is reached. Practically, there will only be a limited research time over which the process can be iterated. In addition, experts' availability should be taken into account, as well as possible fatigue resulting from too many iterations.

Consensus criteria It is important to establish explicit criteria to decide when consensus is reached. For instance, the researcher may establish a threshold, say when a certain percentage of the experts agree^{*}.

* A 75% threshold is often used in the literature

Location The method is usually executed remotely as there is no direct interaction between the experts.

0.2.5.2 Other things to think about

Resources This method is resource intensive, particular if you need to conduct several rounds, so you must ensure commitment from your participants for the whole duration of the study. If you are short of time, a focus group may be a better option.

Lack of discussion The feedback at each iteration is mediated and controlled by the researcher and there is no direct interaction or discussion among the experts. If a deeper investigation of ideas is needed, then other methods should be used, for instance interviews or focus groups.

Difficulty in reaching consensus Consensus may be difficult to reach, particularly, if you are investigating a particularly complex or contentious issue, or you are hoping for predictions concerning highly uncertain or volatile contexts. In such cases, you should reflect of the extent consensus is needed and, if not, then consider alternative methods which may allow you to explore alternative or contrasting views and positions.

0.2.5.3 Further reading

*** find more/better references

Skulmoski, G.J., Hartman, F.T. and Krahn, J., 2007. The Delphi method for graduate research. Journal of Information Technology Education: Research, 6(1), pp.1-21.

0.2.6 Journaling

Journaling is a method requiring participants in a research study to keep regular written personal records in the form of a diary[•] of their experiences and observations during the study. Participants are encouraged to engage in self-reflection in order to surface their inner thoughts, feelings, motivations and perceptions.

Activity: Do I need to know about Journaling?

#14

Check back to your chosen research strategy from Stage 3. Does it involve data generation using Journaling? If so, read through this section and complete the activities.

[•] Or journal, hence the name.

This method generates rich and detailed qualitative data on the participants' subjective experience. It can provide deep insights into complex phenomena, including how things change over time, and allows the

collection of data on everyday experience in naturalistic settings. Because of these characteristics, it is often applied in ethnographic and grounded theory research.

0.2.6.1 Procedural considerations

Diary form Journaling can make use of either hand-written or digital diaries[•]. One or more diaries can be used for different aspects of the journaling process.

- Of course, accessing hand-written notes will be more labour intensive than electronic ones.

Prompts and guidance The goal of the self-reflection should be established at the beginning by the researcher and clearly explained to the participants. Journaling prompts are used to ensure that diary entries align with the objectives of the research, that is to guide participants' focus towards diary entries that will be useful to the research. Prompts can take the form of questions or comments on those diary entries. Lack of guidance or clear instructions is likely to lead to poor data, which inconsistent or incomplete.

Participants Journalling is demanding for participants as maintaining regular diary entries over lengthy periods can be challenging. Therefore, recruiting participants willing to commit to journaling for the duration of the study is essential to the successful generation of comprehensive data. In addition, engagement with journaling should be monitored throughout the study, and perhaps research goals and prompts re-stated to help participants refocus their effort as necessary.

Managing data Journaling can generate large volumes of data from multiple participants, which must be managed carefully, and raises issues of confidentiality, privacy and more generally data protection. As a result it is essential that the researcher establishes a systematic approach to storing and backing up data, whether physical or digital.

0.2.6.2 Other things to think about

Subjectivity and bias Participants are in control of their diary entries, which, as a result are influenced by their emotions, beliefs, preconceptions and cognitive limitations. These, in turn, lead to a number of well recognised biases such as confirmation bias – focusing on evidence in support of prior beliefs, memory and recall biases – difficulties in remembering or reporting accurately past events, or social desirability bias – only making entries deemed socially acceptable or desirable. Awareness of such biases and the use of triangulation may help mitigate resulting weaknesses.

Limited generalisability Diary entries are personal, subjective and usually specific to a particular context or setting. As a result, there may be limited scope for generalising findings based on journaling. If generalisation is an important goal of your research, then you should consider different data generation methods, like questionnaires involving random sampling.

0.2.6.3 Further reading

Activity: Deep dive into journalling

#15

To find out more about journalling, take a look at these resources:

kadarisman2017classroom;
burns2009action;
james2005journaling;
peoples2020write;
feinblum2016journaling;
hayman2012journaling;
mcgrath202115;
taylor2006research;
ovens2020weaving;
giguere2012self-reflective;
bacon2014journaling

0.2.7 Fieldwork

wolcott2005art offers salutary advice:

Quote

There may be discomfort and hardship aplenty connected with the experience, ranging from the distractions of diarrhoea or lost luggage to the despair of personal failure or lost hope, but the extent of one's suffering and sacrifice are not factored into judgments about the worth of the fieldwork as fieldwork.

Fieldwork is a data generation method which requires the researcher to collect data directly in a natural setting[•]. The goal is to gain firsthand knowledge of the phenomena under study, and the method is widely applied across disciplines, including anthropology, sociology, archaeology, geography or environmental science.

Activity: Do I need to know about fieldwork?

#16

Check back to your chosen research strategy from Stage 3. Does it involve data generation using fieldwork? If so, read through this section and complete the activities.

Fieldwork encompasses all kinds of data collection performed directly by the researcher in that setting, be that observations and measurements, collection of samples and specimens, detailed descriptions of direct experience, or other. In sociological studies, fieldwork requires personal involvement of the researcher with the social activities under study, so that participant observations are commonly used.

Through fieldwork the researcher can generate rich, contextualised data to provide deep insights into the phenomena under study. It is particularly suited to situations in which such data cannot be obtained in any other way, and may also lead to unexpected discoveries. It is also flexible in that the researcher can pick data collection techniques to match the specific context and situation.

Fieldwork is not confined to exotic, distant locations[•]! Instead, it can be used in all natural settings, including, say, the researcher's own university or workplace.

- The 'field', hence the name.

- Although fieldwork may have this characteristic for some lucky researchers.

0.2.7.1 Operational considerations

Logistics, equipment and budget Depending on where the field is located, fieldwork may require some detailed planning addressing travel arrangements and accommodation, as well as access to the research site. If in a foreign country, all sort of factors must be considered, including the local transport networks, administrative processes you may need to go through, the climate, etc. You may also need specialised equipment on site, say field equipment and tools for data collection, alongside your personal protection. This can all be quite expensive, so that careful budgeting and securing the required funding in advance is essential.

Permissions and ethical issues If access to the field of interest is restricted, then you will need to gain appropriate permissions to proceed from the relevant authorities. This covers anything from obtaining

permits and licences to access, say, an archeological site, to permission from your employer to perform participant observations in your workplace. The time and effort to obtain such permission must be considered upfront, and all ethical and legal implications factored in. In addition, the study must be conducted in full respect of local culture and norms.

Health and Safety In working in the field, you, and anybody else participating in the research, may be exposed to all sort of hazards, so that assessing and mitigating health and safety risk is paramount. This may involve the introduction of safety protocols, of appropriate training, and appropriate contingencies in the case of an emergency.

Managing data Fieldwork can generate large volumes of data, and may also include precious samples and specimens. All that is collected must be managed carefully, so that alongside issues of confidentiality, privacy and more generally data protection, you may also have to worry about the physical security of those samples and specimens. For this, you will need a systematic approach to storing, protecting and backing up your data, whether physical or digital.

0.2.7.2 Other things to think about

Quality of results The quality of results obtained from fieldwork depends on the data generated in the field, which, in turn, depends upon the skills of the field worker in relation to the specific techniques applied. For instance, using standardised measuring tools will increase the reliability and accuracy of measurements; a reflexive approach will help mitigate against the researcher's personal bias, including confirmation bias; triangulation may mitigate observer and social desirability bias in participant observations, etc. Whichever techniques you choose to apply in your fieldwork, you should ensure you are aware of their potential weaknesses and adopt appropriate strategies to mitigate their effect on the outcomes of your research.

Logistic challenges The logistical challenges to organise fieldwork may well be beyond what can be addressed in the limited time of a Masters project, unless you are able to contribute to a wider research effort, perhaps led by your supervisor, where all logistical issues have already been addressed.

Time and cost Fieldwork can be time consuming both for data collection and analysis, and expensive if travel is required. If time or cost are an issue in your project, then you should consider more time efficient or cheaper alternatives.

0.2.7.3 Further reading

enwiki:1211911661

According to **wolcott2005art**,

Activity: Deep dive into field work #17

To find out more about field work, take a look at these resources:

wolcott2005art;

randall2007fieldwork

is this good to include? Check!

- **wolcott2005art**'s book is both detailed and entertaining. Reading it gives one a feeling that spending time in fieldwork with "Harry" would have been an education in itself. The book includes the importance of laundry to fieldwork, for instance, with experiences of fieldwork in a Canadian Indian reserve to illustrate.

0.2.8 Documents

Existing documents can be used as data sources in order to develop new insights or answer research questions. Research which takes this approach is called document-based research[•]

Activity: Do I need to know about documents? #18

Check back to your chosen research strategy from Stage 3. Does it involve data generation using documents? If so, read through this section and complete the activities.

As a researcher, documents – in the form of academic articles – will already occupy a large proportion of your time/brain/computer. Your collection of academic papers could currently be as many as 50 or more. You will, therefore, already have good experience of interacting with documents and those interactions may well have already helped you gain valuable insights for your project.

Other documents can also be used as data sources in document-based research. The term 'document' is used here in a broad sense to refer to all text-based documents, but also visual materials – such as paintings, maps or photographs, video and audio recordings, and any digitally stored information.

Researchers engage in document-based research by systematically examining and interpreting these documents to extract meaningful information. The documents may be of interest because of their content, or their relation to other documents, or could be studied to discover what they may reveal about their authors, or the historical or cultural context in which they were created. Therefore, a researcher's may have a direct interest in the factual content of a document, or be interested in what that content may indirectly say about

- Or documentary research.

some other phenomena of interest. An example may help clarify the difference between these two modes. This is a copy of a passage from **fynes1873the-miners**:

Quote

Miner :– I believe you have something like 150 collieries to inspect?

Mr. Dunn :– Yes.

Miner :– Twenty-eight in Cumberland?

Mr. Dunn :– Yes.

Miner :– Do you think you are able to inspect all these?

Mr. Dunn :– Well, the Government thinks I am able, you know.

Another Miner :– Were you satisfied with the one shaft at this colliery, if so there is an end to the matter; if not, what steps did you take to remedy the defect? Did you apply to the Secretary of State, showing him that it was defective?

Mr. Dunn :– At this very moment there are three of the largest collieries in Northumberland – Seaton Delaval, North Seaton, and Newsham – managed by the most talented men in Northumberland, all with single shafts. Now, what would you have me to do? Do you think it is my duty to call in question the management of these pits?

Miner :– Am I to understand this is an answer to my question?

Mr. Dunn :– Well, I am not so well satisfied as if they had two, but I have not the power to alter it.

Here, a direct reading could be to identify collieries in which a single shaft existed at that time. An indirect reading could be to explore social relationships within a mining community in 18th century England.

Document-based research provides the researcher with evidence of historical events or social phenomena, including nuanced details and perspectives that may not be available through other means. In particular, documents allow the researcher to access data pertaining to different time periods, locations and cultural contexts. It is a flexible method that can be integrated in several research strategies.

0.2.8.1 Procedural considerations

Accessing documents You need to ensure you have access to the documents you need for your research. While documents are increasingly digitised and easily accessible online, it is also the case that access for many may be restricted by either policies or physical restrictions, say you wish to study restricted confidential

documents in an organisation, or access rare or ancient manuscripts kept in a museum. Ensuring you have the right access at the right time in your project is essential, but can be time consuming, particularly if there are bureaucratic processes you need to go through to gain permission. Related to access are issues of translation if the documents are in a language you are not familiar with, or transcription, if your sources are audio or video recordings. As well as being time consuming, these processes may introduce errors which may be difficult to spot. Finally, it is essential you ensure that your documents are authentic: only using trusted sources is a way to do so.

Selection criteria You must develop clear and explicit selection criteria to choose which documents to include in your study, based on your research problem or question, and aim and objectives. Such criteria should help you collect an appropriate and representative selection of documents for your research, guiding you in what to include and what to leave out. The criteria should ensure that your data sources are diverse and no selection bias creeps in, which may lead to certain positions, perspectives, or types of documents to be either overrepresented or underrepresented.

Data management Alongside generic issues of data storage, protection and privacy, you also need to ensure that your source documents remain easily accessible and that their integrity is maintained: this is both to allow you to revisit those documents repeatedly during your study, and to allow other researchers to check your sources to verify and validate your findings. Whenever possible, you should digitise your source documents to enhance their accessibility and preservation.

0.2.8.2 Other things to think about

Bias in documents Documents are created by people, who necessarily inject their own personal bias into their content[•], which in turn is the result of their historical, cultural and social contexts[•]. In addition, particularly in the case of ancient manuscripts, the documents that have survived may only provide a partial historical account[•]. Being aware of all these biases is essential to the interpretation of documents content and how they may skew or limit the research results.

- So called ‘creator’ bias.
- Another bias, called ‘contextual’ bias.
- You’ll have guessed there is a name for this too, which is ‘survivorship’ bias.

Interpretation challenges Documents are unlikely to provide a complete picture of phenomena under study, partially due to their inherent biases, but also because they may be incomplete or lack crucial details or contextual information may not be available to the researcher. Also certain phenomena may be more documented than others, so that the availability and quality of documents can vary widely across topics,

history or geography. All these factors affect the researcher's ability to interpret their content and draw robust conclusions. It is therefore essential to reflect on the validity and reliability of the documents used in the research, and be open and upfront as to all its potential weaknesses.

Time and effort Document-based research may require large volumes of materials to be selected, collected and analysed, so that it can be very time-consuming. If time is an issue in your project, then alternative data generation methods should be considered.

0.2.8.3 Further reading

Activity: Deep dive into documents

#19

To find out more about documents, take a look at these resources:

coffey2014analysing;

schreier2014qualitative

0.2.9 Computational thinking

Clearly, part of computation thinking is programming a computer to achieve a goal.

However, **wing2006computational**[•], in her **wing2006computational** article that popularised Computational Thinking **wing2006computational** as a candidate curriculum entry, places programming in the broader context of computational thinking. If you are aiming at computational thinking as a way of focussing on your programming skills, it may be that you have misjudged this wider context.

wing2006computational says:

Computational thinking [...] has the following characteristics[•]:

Conceptualizing, not programming. Computer science is not [just] computer programming. Thinking like a computer scientist means more than being able to program a computer. It requires thinking at multiple levels of abstraction;

A way that humans, not computers, think. Computational thinking is a way humans solve problems; it is not trying to get humans to think like computers.

• **wing2006computational** is American, and so uses American spelling – -izi-, artifact, etc.

• **wing2006computational** includes other characteristics, of relevance to curricula but not so much to research.

Complements and combines mathematical and engineering thinking. Computer science inherently draws on mathematical thinking, given that, like all sciences, its formal foundations rest on mathematics. Computer science inherently draws on engineering thinking, given that we build systems that interact with the real world.

Ideas, not artifacts. It's not just the software and hardware artifacts we produce that will be physically present everywhere and touch our lives all the time, it will be the computational concepts we use to approach and solve problems, manage our daily lives, and communicate and interact with other people.

Recognising the social context of computational thinking is, thus, critical to research in the area.

Activity: Do I need to know about computation thinking? #20

Check back to your chosen research strategy from Stage 3. Does it involve data generation using computation thinking? If so, read through this section and complete the activities.

The techniques generally associated with computational thinking are:

- Decomposition: breaking a problem into parts that are easier to solve;
- Pattern recognition & Generalisation: seeking significance from repeated structures within data
- Abstraction: building higher level representations of things
- Algorithms: associating behaviours with computational structures
- Evaluation: called testing for software, but also has wider applicability when an algorithm is socialised, i.e., used within its final social context.

0.2.9.1 Computational thinking tools

Given the explosive growth in the use of computers over the past half century, you may not be surprised to hear that there are thousands of useful^{*} computer tools to support computational thinking. Fortunately, for the vast majority of these, they support one or more computational models, which are the bases of computational thinking. Those computational models fit into the following categories

^{*} As well as some that are less than useful!

- computational mode: sequential, concurrent, agent based. Example tools: flowcharts, UML, Petri Nets, •
 - computational paradigm: imperative, declarative. Examples: procedural, object oriented; logical, functional, reactive; etc.
 - language: machine, low-level, high level, interpreted: machine code, C, Swift, Python.
 - developmental process: lying on a continuum between plan-driven and micro-iterative
 - requirements form: functional, non-functional;
 - data structure: database, spreadsheet, json, etc.
 - delivery platform: web, package, snippet, 
 - maintenance mechanism: GitHub,
 - language libraries: React, Ruby on Rails, etc. •
- [more here](#)
- Computational modes tend to be packaged for use, so – although the basis of a computational mode might be concurrent, you would not have to deal with the intricacies of modelling in, say, Petri nets. Unless, of course, your intention was to investigate their properties.
 - Language libraries tend to come into and go out of date regularly as new paradigms of delivery are created and superseded.

0.2.9.2 Workflow

The main idea behind computational thinking is to link a wished-for behaviour to a computational structure. Here, we're thinking of the wished for behaviours of a client in context being translated to an algorithm, data model, or other computational element.

The ways in which the wished-for behaviour of a client in context are understood is through the discipline of requirements analysis. Although naive requirements analysis is possible for simple projects – asking a client what functional behaviours they need – it is more often driven by the explicit or implicit management or developmental risk, the risks being those of

- misunderstanding the context of deployment – where the final system will be installed
- misunderstanding the problem that is to be solved – what functional and non-functional behaviour is wished-for;
- misunderstanding the solution space – what technologies can be employed within the context and how they should integrate with human elements of the system

- misunderstanding of the change path by which any current system will become the new system
- misunderstanding the process needs for developing and delivering the system
- [more here]

0.2.9.3 Other things to think about

If you don't have experience with programming languages, then you can learn, but the learning curve can be very steep[•].

Engaging with clients as stakeholders in a development can also be challenging, in that it makes great demands on both parties' time – there can be a great deal of iteration needed to get to a form where the client needs can be satisfied. Add to this validation of any delivered product by the client and a need for good project management skills come to the fore.

Often, developmental research projects constitute the research component of professional body-accredited degrees. Such degrees add requirements for good development practice including, for instance, documentation of the development process and its inputs, testing regimes, etc. These requirements can add substantially to the basic task of creating a solution, but aim to substantiate real-world problem solving.

[More here]

Activity: Deep dive into computation thinking

#21

To find out more about computation thinking, take a look at these resources:

[angevine2017computational](#);

[figueiredo2017improve](#);

[lyon2020computational](#)

- This video will give you some idea of what will be required: <https://www.youtube.com/watch?v=fLMZAHyrypo>. Stephen Wolfram is a particularly interesting, if restless, speaker and a leader in this area.

But this isn't a great video, nor is any other that I've looked at, as they're created by computer scientists... They're also heavily advertised.

0.2.10 Mathematical thinking

Mathematical thinking has had many centuries more than computational thinking to develop and the tools that exist as part of it are very stable. They are also much better explored and – due to the efforts of many great mathematicians, are more complex to apply to achieve their full potential.

Having said that, mathematical thinking has less of a basis in real-world problem solving, although mathematical techniques can apply to real-world problems, the difficulty being that closed form solutions, as mathematics tends to create[•]. Mathematics has limits of applicability, some of which are extremely subtle, even in relatively simple situations: complete closed form solutions in radicals do not exist for problems as simple as finding the roots to polynomial equations of order 5[•] or above; the general three body problem, for instance, being resistant to differential equation analysis; and others.

Alternatively, so called *numerical approximations* can be made to most problems, but these are often related to computational solutions and so might be better approached through computational thinking.

Another alternative that again takes us back to computational thinking, is the use of agent based simulations in which concurrently acting independent agents are used.

Activity: Do I need to know about mathematical thinking?

#22

Check back to your chosen research strategy from Stage 3. Does it involve data generation using mathematical thinking? If so, read through this section and complete the activities.

[More here.](#)[\[More here\]](#)

0.2.10.1 Mathematical thinking tools

Differential equations; matrices; algebra; numbers theory; sets; functions; relations; logics[•]; topology; geometry; calculus; algebra; analysis;

According to **burton1984mathematical**, there are four processes that are central to mathematical thinking:

Specialising exploring a problem through examples. Each example provides the opportunity for manipulating elements that are concrete, whether they are physical manifestations or ideas.

Conjecturing when enough such examples have been examined, you can conjecture about the relationships that connect them. Through conjecturing, underlying patterns are explored, expressed, and then substantiated.

Generalising if you are lucky enough to have found a pattern, then you might try to generalise it to creating order and meaning out of a – potentially, overwhelming – amount of data.

- A *closed form* solution is something like a set of differential equations that completely describe a – typically idealised – problem.

- So called *quintic equations* ([enwiki:1209364767](#)). A radical is a mathematical expression involving only the coefficients of the equation, and the basic arithmetic operations (addition, subtraction, multiplication, division, and taking the nth-root).

[More here.](#)[\[More here\]](#)

- *Logics* is plural as there are many, depending on which area of mathematics you are using.

[More here?](#)

Convincing a generalisation must be tested until it is convincing to the reader – this is the basis of the knowledge contribution from mathematical thinking.

0.2.10.2 Workflow

Working from simple examples to more complex through an iterative workflow that establishes an appropriate mathematical approach. Building interpretations of a problem which are amenable to existing mathematical tools, or which suggest extensions to them. Developing and applying those tools to produce results. Iteration towards increasingly complete mathematical solutions, refinement of the problem to its important mathematics constituents. Formalisation and – if necessary – proof of results. Application.

0.2.10.3 Other things to think about

Mathematical thinking is arrived at through creative thinking and deep study of mathematical tools and techniques. The sophistication of mathematics often means that, either:

- a particular area of research has already been taken past the abilities of a masters-research-level mathematician;
- it is not amenable to (current) mathematical tools and techniques, and further creative[•] mathematical thinking will be necessary to progress.
- And, most likely, deep and advanced, out of the box, out of this world, and further.
- What is often missing from the mathematical literature – or what isn't always visible to the new entrant – is the often vast timescales over which mathematical progress is made. Bertrand Russell and Alfred Whitehead spent over two decades of their professional lives in the creation of the three volume *Principia Mathematica*. A fourth volume – on geometry – was begun but never completed. In another example, the proof of Fermat's Last Theorem took 358 years to complete.

Although neither of these characteristics are insuperable, they make timely contributions to knowledge through the application of mathematical approaches difficult[•]. It's worth moderating your expectations of what can be achieved in mathematical research at masters level – discussion with your supervisor of what their expectations are would be very worthwhile.

Activity: Understanding your supervisor's mathematical thinking expectations #23

Schedule some time with your supervisor to discuss what they hope will be achieved through your research project.

Mathematics abstracts from real-world complexity: in modelling traffic to improve flow through a complex junction, for instance, one would not necessarily consider the economy of individual cars, or the noise pollution created by a solution. This can reduce a real-world problem to a complexity that is approachable, but may

also lead to non-solutions when applied back in the real world, for instance, leading to complaints from local home owners that noise pollution has risen through a solution.

Mathematics is a very tight community with very high publication standards. The language of mathematics is dense[•].

Because of this, many mathematical research projects at masters level are designed to provide a way into the mathematical literature. A supervisor will set a mathematical task that may already have been solved. The contribution a research student might make, then, is not a contribution to knowledge in the formal sense of extending mathematics, but the widening of the mathematical community to include another researcher whose skills have expanded to be able to make a novel restatement of a problem, for instance, and research further.

The study of the development of mathematical thinking, for instance, in schoolchildren is a very fruitful area with much still to be contributed.

Activity: Deep dive into mathematical thinking

#24

To find out more about mathematical thinking, take a look at these resources:

[stacey1982thinking](#)

[More here](#)

0.2.11 Statistical thinking

Statistical thinking involves designing a study to collect data, analyse patterns in the data, and draw conclusions that go beyond the observed data. Sampling is key to being able to generalise results, while random assignment is key for cause-and-effect conclusions. Probability models help assess random variation and estimate margin of error.

Activity: Do I need to know about statistical thinking?

#25

Check back to your chosen research strategy from Stage 3. Does it involve data generation using statistical thinking? If so, read through this section and complete the activities.

Source: [chance2024statistical](#).

- A mathematical statement indicative of this complexity is: “Let q be $x^5 - x - 1$. Let G be its Galois group, which acts faithfully on the set of complex roots of q . Numbering the roots lets one identify G with a subgroup of the symmetric group S_5 . Since $q \bmod 2$ factors as $(x^2 + x + 1)(x^3 + x^2 + 1)$ in $\mathbb{F}_2[x]...$ ”.

0.2.11.1 Workflow

0.2.11.2 Other things to think about

Add more definitive resources here

Activity: Deep dive into statistical thinking

#26

To find out more about statistical thinking, take a look at these resources:

[chance2024statistical](#)

0.3 Managing your raw data

Your chosen research strategy may require you to generate data of many kinds and from many sources. The amount of data you collect could be a few ideas or a mound of documents, from a few to terabytes of data.

Before proceeding with your data analysis, you must ensure your data are properly organised and stored, so that you don't lose track of important information, and you can easily locate and refer back to appropriate data during your analysis and when writing up your research – you might need to include a representative sample of your raw data in an appendix of your dissertation as evidence of your data generation, for instance, identifying which sample can take some time if done afterward while it can be immediate if done while.

Although techniques for doing so are outside of the scope of this book, your data storage should be secured against data loss either due to technical issues, such as computer failure, or due to a data breach, such as through hackers, at least to the standards required by law, any additional requirements made by your organisation, those of any participants, their organisations, and any other stakeholders[•].

It is also important that you put your raw data in a form which is useable for analysis. Spreadsheets are particularly useful for this purpose, especially if your data is quantitative, so that this is a common way to organise and store raw data. In fact, most publicly available data sets used in research and beyond are stored as spreadsheet files: if you are going to use one such data set, then your raw data are likely already organised for you!

Spreadsheets organise data in rows and columns, so that you can easily enter your raw data using rows for your observations/measurements and columns for your variables. As we will see later on, spreadsheets come a wide range of functionalities for data manipulation and for some level of data analysis. They are also easily extensible, so that you can grow your data sets incrementally.

- Once upon a time, in a galaxy far, far away, data generation and storage used to be a *laissez-faire* thing. Today, your organisation can be fined vast amounts of money for any data misuse, so they tend to take it more seriously. If data loss were to happen, amongst other things, it'd probably mean you'll fail your degree.

Activity: Organising and storing your raw data

#27

Consider the data and evidence you have collected or are planning to collect. List the actions you have taken/will need to take in relation to:

- organising your raw data
- storing and backing up your data
- protecting personal data

Make sure you complete those actions before moving on to your data analysis.

With your raw data properly organised and stored, you can now proceed to your data analysis.

0.4 Analysis methods

Your choice of data analysis methods is part of your research design, and relates to the kind of data and evidence you have generated, and what you are trying to achieve, that is your aim and objectives.

This section provides an introduction to some common analysis methods. It is far from complete and does not go very deeply into the details of each method: entire books have been written on any of them! By studying this section, you won't become an expert in any of these methods, but you will have gained enough understanding to be able to make a judicious selection for your project. After that, you should review the related specialised literature to help you apply your chosen methods appropriately. You should also talk regularly to your supervisor for further guidance.

0.4.1 Using tables to analyse data

The following kinds of tables are used extensively in research and often found in dissertations.

Pivot tables Pivot tables can be used to summarise, sort, filter, re-organise or group data organised in rows and columns, and perform calculations on them, such as counting, generating totals or averages, and much more. Pivot tables are both powerful and versatile[•], and one of the most widespread tools for data

• In fact, they are so versatile that we'll only be able to provide few illustrative examples. Much, much more can be found online!

analysis.

You can generate a pivot table from any data set organised in rows and columns, regardless of whether the values are quantitative or qualitative: all common spreadsheet applications[•] include this function.

figure 0.1 gives an example: these are the first few raws of a data set related to the US housing market[•]. The dataset contains over 9,316 entries, each corresponding to a distinct property. Each property is characterised by a number of attributes: size in square feet, number of bedrooms and bathrooms, type of neighbourhood, the year it was built and its market price in US dollars. As you can see, this table includes both numerical and categorical variables.

ID	SquareFeet	Bedrooms	Bathrooms	Neighborhood	YearBuilt	Price
1	2126	4	1	Rural	1969	US\$ 215,355.28
2	2459	3	2	Rural	1980	US\$ 195,014.22
3	1860	2	1	Suburb	1970	US\$ 306,891.01
4	2294	2	1	Urban	1996	US\$ 206,786.79
5	2130	5	2	Suburb	2001	US\$ 272,436.24
6	2095	2	3	Suburb	2020	US\$ 198,208.80
7	2724	2	1	Suburb	1993	US\$ 343,429.32
8	2044	4	3	Rural	1957	US\$ 184,992.32
9	2638	4	3	Urban	1959	US\$ 377,998.59

Figure 0.1: first few rows of the example dataset

Pivot tables can be used to summarise such data to answer certain questions. For instance, we may be interested in the average house price by neighbourhood and number of bedrooms, which would result in the pivot table in figure 0.2, which gives the average price of each combination. The ‘grand totals’ in the table are also averages, by row and by column.

Alternatively, we may be interested in finding out how many properties of each kind have been built in each neighbourhood. In this case the pivot table would look like that in figure 0.3. The grand totals in this

- From MS Excel to Apple Numbers to Google Sheets.
- It was taken from one of Kaggle’s free datasets, the housing price dataset. Kaggle is possibly the largest and best known online community for data science and machine learning.

Neighborhood	Rural	Suburb	Urban	Grand Total
Bedrooms	Price (Average)			
2	US\$ 218,323.92	US\$ 216,300.13	US\$ 220,050.01	US\$ 218,230.99
3	US\$ 219,053.37	US\$ 220,397.86	US\$ 223,737.67	US\$ 221,057.88
4	US\$ 227,774.55	US\$ 224,609.50	US\$ 230,086.56	US\$ 227,473.37
5	US\$ 231,112.60	US\$ 231,776.73	US\$ 234,894.98	US\$ 232,595.48
Grand Total	US\$224096.13	US\$223234.19	US\$227166.20	US\$224827.33

Figure 0.2: Pivot table of average property prices by neighbourhood and number of bedrooms

case are counts. Note how we have added combinations of bedroom and bathroom numbers to characterise each type of property.

These are just but two examples of questions about the data you can address by using pivot tables, out of a vast range of the possibilities. If your data are organised in tables, then it is well worth spending some time becoming familiar with pivot tables.

Activity: Pivot tables in Excel

#28

Download the housing price data set from Kaggle and re-create the pivot tables in our example. Come up with other questions you could ask of the data and generate related pivot tables.

Guidance

Feel free to use your preferred spreadsheet application for this activity, as long as it supports pivot tables – most do.

You may have to register with Kaggle to gain access to the data set.

The Excel Help facility and documentation provides all the info you need to create a pivot table. However, you could also browse some of the very many freely available online resources and tutorials on this topic.

Neighborhood		Rural	Suburb	Urban	Grand Total
Bedrooms	Bathrooms	ID (Sum)			
▼ 2	1	180	220	261	661
	2	129	214	283	626
	3	155	45	75	275
2 Total		464	479	619	1562
▼ 3	1	96	130	129	355
	2	397	149	245	791
	3	127	139	889	1155
3 Total		620	418	1263	2301
▼ 4	1	142	306	228	676
	2	453	379	410	1242
	3	315	175	194	684
4 Total		910	860	832	2602
▼ 5	1	325	282	378	985
	2	438	127	10	575
	3	534	367	390	1291
5 Total		1297	776	778	2851
Grand Total		3291	2533	3492	9316

Figure 0.3: Pivot table of property counts by neighbourhood and number of bedrooms/bathrooms

Frequency and contingency tables Frequency tables are used to summarise the frequency (or count) of values taken by a categorical variables in a data set. For instance, after studying a degree, a student's outcome may be classed as distinction, merit, pass or fail. A frequency table can then be used to summarise the frequency of each class of outcome for a particular students' cohort, as shown in Table 0.2.

Table 0.2: Example of frequency table

	Distinction	Merit	Pass	Fail
Outcome	12	26	42	5

Contingency tables[•] are a form of frequency tables used to tabulate the value frequencies of two categorical variables. For instance, following from our previous example, we may like to tabulate the outcome value frequencies in the cohort against gender, as shown in Table 0.3.

• Also known as *cross-tabulation* tables.

Table 0.3: Example of contingency table

Outcome by Gender	Distinction	Merit	Pass	Fail
Female	7	12	21	2
Male	5	13	19	3
Other	0	1	2	0
Totals	12	26	42	5

Contingency tables are frequently used to summarise and analyse data collected in survey research, and are a key tool in statistical analysis.

Both frequency and contingency tables can be generated as pivot tables in a spreadsheet. In fact, the table in figure 0.3 is a contingency table.

0.4.2 Statistical analysis

Statistical analysis is an umbrella terms for a set of methods which can be applied to numerical and categorical data. More precisely, in statistics data types are classified as:

- scalar, which includes all measurements and counts; with reference to the types in Section ??, these are all numerical data, continuous, discrete, interval and ratio data.
- categorical, both ordinal and nominal.

There are two broad categories of statistical methods:

- descriptive statistics, whose aim is to describe data; and
- inferential statistics, whose aim is to make predictions from data.

We briefly consider each in what follows.

0.4.2.1 Descriptive statistics

These are used to describe various attributes of a data set. The basics are:

- count, to establish how many entries there are in the data set
 - centrality, to establish the ‘centre’ of the data set. Three measures are commonly used: the *mean*, which provides the average value of the data set; the *median*, which provides its mid point[•]; and the *mode*, which indicates the value that occurs most frequently, if any[•].
 - dispersion, to establish the spread of the data in the data set. Range and standard deviation are two common measures. The *range* is the difference between smallest (minimum) and largest (maximum) values. The *standard deviation* is based on a mathematical formula which considers the distance of each value in the data set from the mean. It is not essential for you to know such formula, which is automatically computed by spreadsheets and statistical software[•]. The larger the standard deviation, the greater the dispersion.
 - skewness, to establish how symmetrically distributed the values in the data set are in relation to the centre. In the case of perfect symmetry, skewness is equal to zero, and mean and median are equal. When asymmetric, mean and median are different and the distribution may be either right (mean smaller than median, and negative skewness) or left (mean greater than median, and positive skewness) skewed. A perfectly symmetric distribution is usually referred to as a *normal distribution* or *bell curve*, from the shape of the line that can be obtained by plotting the data on a chart[•].
- Remember that quantitative data can be ordered.
 - There is no mode if no value is repeated in the data set.
 - Of course, you can always look it up in the literature...
 - This oversimplifies the topic in order to give some intuition in case you have not come across these terms before. A lot more should be said about the normal distribution and its pivotal role in statistics!

Not all descriptive statistics apply to categorical data. In particular, the mode is used as the main measure of centrality for nominal data, while the median is used for ordinal data which are not numeric.

These are lots of definitions to digest, particularly if you haven't encountered these terms before! The following activity should help.

Activity: Descriptive statistics in Excel

#29

Assume you have measured the weight in grams of each apple in a basket, obtaining the following numbers: 105, 120, 122, 125, 127, 128, 129, 130, 132, 133, 135, 135, 138, 140, 128. Enter these data in an Excel sheet and use its built-in data analysis function to generate the related descriptive statistics.

Guidance

In the current version of Excel, you can access this function from the Data tab, by pressing the Data Analysis button. If you find it difficult to locate this function, you should refer to the documentation or to some of the many tutorials on this topic which are freely available online.

Discussion

You should have obtained the following values:

Attribute	Value
mean	128.47
median	129
mode	128
standard deviation	8.55
skewness	-1.4
range	35
minimum	105
maximum	140
count	15

There are 15 values in this data set, with range 35 (the difference between maximum and minimum values). In terms of centrality, the mean (128.47) is slightly smaller than the median (129), and Excel reports a mode at 128. In reality, if you look at the data you will see that there are two modes in this data set^a, 128 and 135, but Excel only returns the first encountered! In terms of dispersion, the standard

deviation is telling us that most apple weights are within 8.55 grams of the mean (below or above), so the apple weights are similar in the apple baskets. Note that the skewness is negative, which is consistent with the mean being smaller than median, so the data distribution is right skewed.

^a Statisticians call this *bi-modal*.

In your dissertation, you can easily present such descriptive statistics as a table, possibly adapting that automatically generated by your spreadsheet.

In addition, charts can be used to visualise the data and examine their descriptive statistics.

With scalar data, like in our example, you can use a *histogram*. The one in figure 0.4 uses the apple weights from the previous activity: on the horizontal axis, we have the distinct weights, and on the vertical axis, their frequencies, that is how many times each weight appears in the data set. Given the values you have obtained for the data set descriptive statistics, you can easily locate on the chart min and max values, and mean, median and mode. In this case, the two ‘peaks’ correspond to the two modes we mentioned in the activity. You can also check that most of the values are within 8.55 grams from the mean, either way: the only values left out are 105 (to the left) and 138 and 140 (to the right). Skewness is not obviously notable on this chart, so that we will use a different chart for that purpose.

Before we do that, however, it is worth noticing that given our small data set of discrete values, we have used a histogram with ‘bin’ size equal to one, which allows us to plot each individual apple weight. A *bin* in a histogram is essentially a way to group a number of values, with bin size establishing the spread of each bin. Frequencies are then calculated by bin. Grouping values in bins is necessary with large data sets and/or with continuous data. figure 0.5 illustrates a histogram for our example, in which the bin size is 5: that is, each bin spans a set of 5 possible values.

In order to visualise both spread and skewness, a useful chart is the boxplot, illustrated in figure 0.6 for our example. This is made of a ‘box’ around the median of the data, and some ‘whiskers’ on each side of the box^b. It is obtained by dividing up the data into quartiles, each containing a quarter (or 25%) of the data, with the median in the centre. The box includes the two quartiles on each side of the median, which, together, account for half of the values in the data set. The whiskers account for the two other quartiles, with a caveat: if there are very extreme values, these are treated as possible outliers and left out of the whiskers. This is, in fact, the case in our example where value 105 is treated as an outlier in the chart: it is a dot on its own, not included in the left whisker. The whisker length provides an indication of spread: the longer the whiskers, the more spread out the data. Instead, the position of the median in relation to

^b Which is why this chart is also called a *box and whiskers* plot.

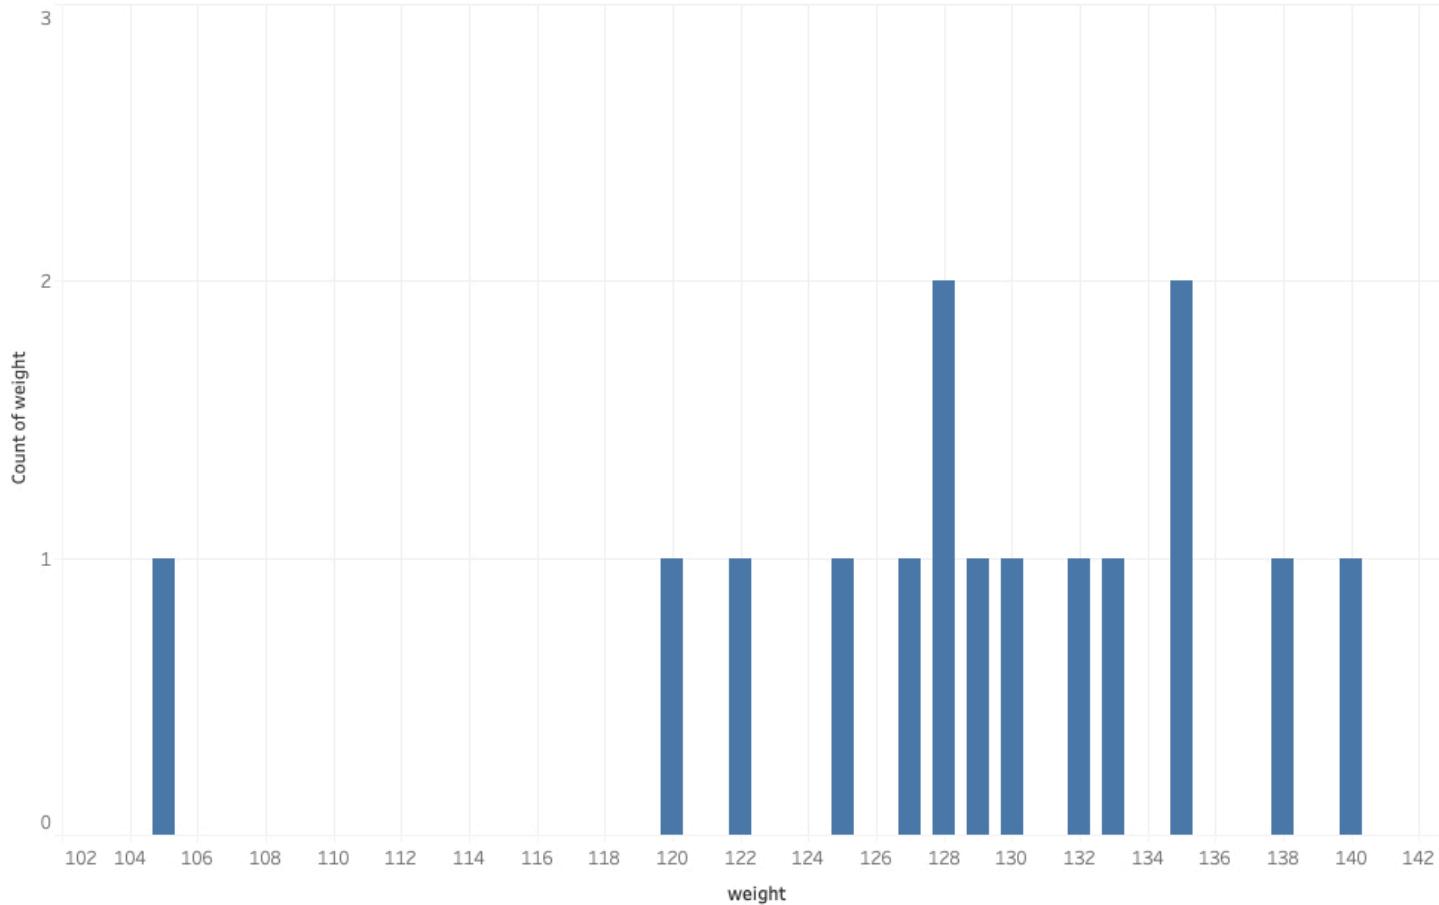


Figure 0.4: Histogram for the apple weights (bin size = 1)

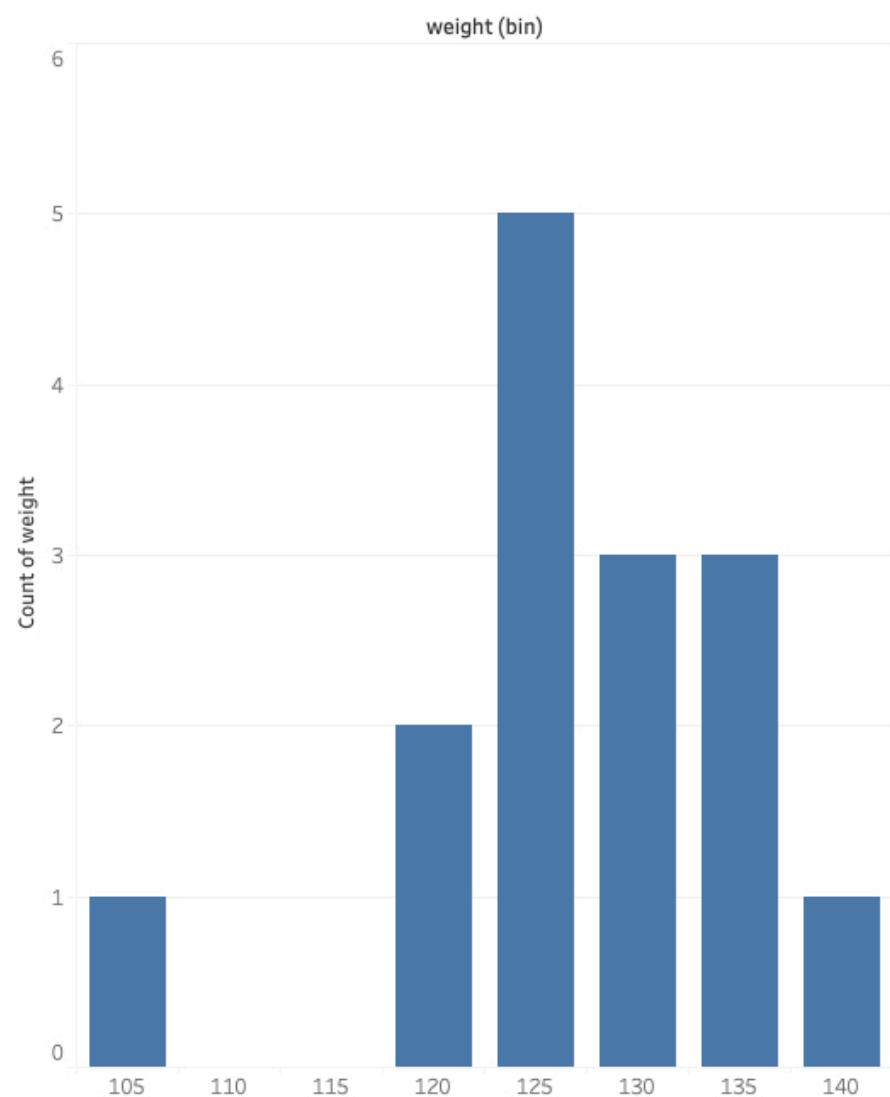


Figure 0.5: Histogram for the apple weights (bin size = 5)

the extreme of the box provides an indication of skewness: in our example the median is further away from the right edge (just!), indicating that the data distribution is slightly right-skewed (consistent with the negative skewness value in the descriptive statistics).



Figure 0.6: Boxplot for the apple weights

To be more precise, the relation between a boxplot and its underlying statistical features is illustrated in figure 0.7. The two quartiles around the median represent the interquartile range (IQR) of the data set. The whisker lengths, calculated based on the formulae in the figure, allows the identification of lower and upper bounds beyond which values are seen as extreme and represented separately as outliers. An outlier, therefore, is just a value which is distant from most of the other values in the data set: it may point to an error, which should be corrected, or an anomaly, which may require further investigation, but that's not necessarily the case. However, it's good practice to investigate all outliers to understand why they have occurred.

Activity: Charts in Excel

#30

Go back to your Excel sheet from the previous activities and generate charts similar to those in the figures above.

Guidance

In the current version of Excel, you can generate these charts from the Insert tab, by choosing from the Statistical charts menu. If you find it difficult to locate this function, you should refer to the documentation or check some of the many tutorials on this topic which are freely available online.

Table 0.4 summarises useful charts that can be applied to visualise your data and their descriptive statistics: it includes charts we have not used in our examples, but which are very common, so that you can find plenty of study materials online should you wish to look them up and use them.

Calculating descriptive statistics and visualising data in appropriate charts, should be the first step in your statistical data analysis, as these provide useful summaries and visualisations of key properties of your




Figure 0.7: The features of a boxplot — LR to redraw as taken from the web

Table 0.4: Common charts to visualise data sets and their descriptive statistics

Chart	Variable(s)	Purpose
bar chart	one categorical	to visualise counts/frequencies/proportions/percentages
staked bar chart	two categorical	to compare counts/frequencies/proportions/percentages between two groups
histogram	one scalar	to visualise distribution, including centrality, dispersion and skewness
scatter diagram	two scalar	to visualise relationships and possible outliers
boxplot	one scale or one categorical	to visualise spread, skewness, median, IQR and possible outliers
line chart	one scalar by time	to visualise change over time

data set. And, as you have found out in the activities, you do not need to be a statistician to be able to generate them!

Descriptive statistics may also help you identify errors or anomalies in the data, and can inform possible follow-up analysis, including inferential statistical analysis. Depending on your research aim and objectives, they could also be all you need in your project.

If you have collected scalar or categorical data, it is time for you to have a go at analysing them using descriptive statistics and charts.

Activity: Applying descriptive statistics to your data

#31

Calculate descriptive statistics for your data set, and generate appropriate charts.

Guidance

MS Excel is relatively straightforward to use for this purpose, but feel free to use other tools you may be already familiar with, including statistical or data analytics packages. Whichever tool you use, you should ensure it supports the functionalities we have discussed in this section.

0.4.2.2 Inferential statistics

Inferential statistics relies on the concepts of population and sample: the *population* is the entire group you are interested in studying – say, all UK voters in a general election; while the *sample* is the portion or subset of that group you have access to in your research. Then the aim of inferential statistics is to establish whether patterns or effects you have observed in the sample can be generalised to, i.e., inferred for, the whole population, or whether they are the result of chance. In inferential statistics this is achieved through statistical tests.

A statistical test tells you whether the proposition[•] you wish to test on your sample is likely to be true in the population under study. For this to work, your sample must be representative of the population[•].

A statistical test returns a measure of *statistical significance*, which is used to provide evidence (or otherwise) that the pattern or effect you see in your sample is also likely to exist in the population, and is not just the effect of chance[•]. As a corollary, if your sample is very large, almost all effects observed in the sample will be likely present in the population; vice-versa, if your sample is very small, most effects observed in the sample are unlikely to be present in the population, unless they are really very large. As a rule of

- You can think of a proposition as an educated guess you have made based on some observations, but that has yet to be supported by evidence.
- We discussed sampling in Section ??
- Contrary to the common language meaning of ‘significance’ as big or important, statistical significance only indicates that the effect is likely to exist in the population, where it may well be small or unimportant!

thumb, most tests require a sample size of at least 30 observations, but more precise sample size estimates can be made based on population size and expected significance level[•].

Each statistical test comprises the following elements.

- **Hypotheses** There are two, *null* and *alternative* hypotheses. Inferential statistics assumes you can't prove something to be true, but you can disprove something by finding an exception. Here is a classic example: you can't prove that all swans are white, but you can disprove they are by finding a black swan! So, you must set the null hypothesis to what you want to disprove about the population, with the alternative hypothesis being what you are really interested in finding out. So, the null hypothesis is usually a statement of no pattern/effect in the population.
- **Significance** This is the level of statistical significance for the test. It's known as the *alpha* (α) value from the Greek name of the mathematical variable used to express it. Most tests are run with $\alpha = 0.05$, which gives a 5% probability that we may infer that the null hypothesis is disproved while in actually it is correct[•].
- **Sample(s)** You need to have one or more (representative) samples of the population of interest on which to perform the test. Multiple samples are used in some tests, typically to compare specific statistics in different groups within the population or changes within a group over time or after an intervention of interest, say treating patients with a new pharmacological drug.
- **p-value** This is the probability calculated for your test by your statistical package, and which is used to decide the outcome of the test.
- **Decision** This is based on the p-value in relation to the α value: if the p-value is less than the α value, then the null hypothesis is *rejected*, i.e. disproved, which means your alternative hypothesis that there is an effect in the population is supported by statistical evidence.

There are very many statistical tests to choose from, depending on the kind of data you have and their distribution, the purpose of your analysis and the number of samples involved.

Statistical tests are applicable to both scalar and categorical data and can be used to compare values of specific statistics or to establish statistical relationships between variables, specifically:

- an *association* between variables means that one variable can be used to provide some information about the other

- Formulae for the ideal sample size are easily found in the literature and online.

- This is called a Type I error in Statistics.

- a *correlation* is a particular type of association such that the two associated variables always change together, for instance they both increase or decrease at the same time, or when one increases the other always decreases.

Statistical tests can be used to estimate the strength of an association (i.e., the extent changes in one correspond to changes in the other) and its direction (whether the variable changes are in the same or opposite way).

We will not detail all possible statistical tests in this introductory section — once again, entire books have been written about them! Instead, we provide Tables 0.5 and 0.6 as summaries of the most common tests that you can then follow up in the literature, should you wish to apply any in your research.

Even if these tests are only a sub-set of all statistical tests available, there is a lot to digest. The next activity should help you use these tables to choose an appropriate test.

Activity: Choosing an appropriate test

#32

Consider the following scenarios: for each, use the information in the tables to decide which test to apply and what the null hypothesis should be. For each, write down your reasoning, choice and null hypothesis.

- *Scenario 1* to investigate the amount of sugar contained in baby food of a particular brand against a recommended threshold, from a sample of 30 products of that brand.
- *Scenario 2* to investigate the number of products per hour of two manufacturing machines in the same plant, by observing the two machines' output over 24 hours.
- *Scenario 3* to investigate the effect of temperature on the consumption of ice cream in a particular city over 12 months.
- *Scenario 4* to investigate whether taste in chocolate types, say white vs milk vs dark, is related to gender in a particular country.

Guidance

To simplify things, always assume normal distributions.

Discussion

Table 0.5: Common statistical tests for comparison. *Parametric* tests apply to normally distributed data (see Section 0.4.2.1), while *non parametric* tests to skewed distributions.

Purpose	Variables	Example	Para- met- ric	Non- para- metric	Notes
to compare the sample mean against a specific value	one scalar	to investigate whether AA batteries of a particular brand have the claimed lifespan	one sample t-test	n/a	
to compare the sample proportion against a specific value	one categorical	to investigate the proportion of people who voted for a particular party in a city against that for the whole country	one sample z-test	n/a	
to compare the means of two independent samples	scalar	to compare the mean scores (dependent) of students studying the same subject with two different teaching approaches (explanatory)	independent t-test	Mann-Whitney test/Wilcoxon rank sum	two samples are <i>independent</i> when there is no reason to believe that observations in one sample are influenced or determined by those in the other
to compare the means of three or more independent samples	scalar dependent; nominal explanatory	to compare the mean scores (dependent variable) of students studying the same subject with three or more different teaching approaches (explanatory variable)	one-way ANOVA	Kruskal-Wallis test	
to compare the average difference between paired samples against a particular value	scalar dependent; time or condition as explanatory	to compare the blood pressure readings (dependent variable) of a group of people before and after exercising (explanatory variable)	paired t-test	Wilcoxon signed rank test	in paired samples each data point in one sample is uniquely matched to a data point in the other sample; this happens, for instance, when we measure a factor before and after an intervention, or take different readings for the same group of individuals. Because of this, paired samples are not independent.

Table 0.6: Common statistical tests for association. *Parametric* tests apply to normally distributed data, while *non parametric* tests to skewed distributions.

Purpose	Variables	Example	Para-met-ric	Non para-metric	Notes
to investigate correlation between two continuous variables	scalar dependent and explanatory	to investigate the relation between blood pressure (dependent) and age (explanatory)	Pearson's Correlation Coefficient	Spearman's Correlation Coefficient	
to investigate association between two categorical variables	categorical dependent and explanatory	to find out if there are gender (categorical) differences in the choice of modes of transport (categorical) in a city	chi-squared	n/a	
to investigate association between two categorical variables when the sample is small	categorical dependent and explanatory	to find out if there are gender (categorical) differences in the choice of modes of transport (categorical) in a city	fisher's Exact test	n/a	the sample size n should be less than 20
to predict the value of one variable from that of one or more other variables	scalar dependent and any kind of explanatory	to predict house prices (dependent) based on location (explanatory, categorical) and number of bedrooms (explanatory, scalar)	linear regression	n/a	linear regression relies on associations between dependent and explanatory variables
to predict the value of a binary variable from that of two or more other variables	binary categorical dependent and any kind of explanatory	to predict whether a customer is likely or not to purchase a certain product (dependent) based on previous purchased products (explanatory, categorical) and average annual spent (explanatory, scalar)	logistic regression	n/a	a binary variable has only two possible values, so that logistic regression calculates the probability of each value based on the values of the explanatory variables. Because of this logistic regression can be used as a classification method

Assuming normal distributions, for each scenario, we have considered:

- the kind of data
- number of samples and their size
- purpose of the investigation

This is what we have concluded:

- *Scenario 1* scalar variable (amount of sugar); one sample of 30 products; to compare the sample mean against the recommended threshold. The test to use is a t-test with null hypothesis that the sample mean is above the threshold.
- *Scenario 2* scalar dependent (number of products per hour) and categorical explanatory (which machine); two samples (one per machine over the time span); to compare the means of products per hours for the two machines; there is no reason to think that the working of one machine may influence that of the other. The test to use is an independent t-test with null hypothesis that the two sample means are different.
- *Scenario 3* scalar dependent (level of ice cream consumption) and scalar explanatory (temperature); one sample over the period; to investigate any relationship between the two variables. The test to apply is Pearson's correlation with null hypothesis that there is no association between the two variables. If, in addition, we wanted to make predictions on ice cream consumption based on temperature, then we could also apply linear regression.
- *Scenario 4* both dependent (chocolate taste) and explanatory (gender) are categorical; one sample from the country; to investigate association. The test to use is a Chi-squared with null hypothesis that gender has no association with chocolate taste.

0.4.3 Qualitative analysis

to cover the need to establish an analytical framework to help apply these techniques; also include the various definition of content/narrative analysis etc

As indicated in Section ??, common methods for qualitative data analysis include content, thematic, narrative or discourse analysis. While their goal may be different, they all apply the initial step of *coding*, which we discuss next.

0.4.3.1 Coding qualitative data

A *code* is essentially a label which describes an extract from qualitative data set, with *coding* the process of creating and assigning codes to categorise those extracts.

Coding is important and it helps you ensure that your analysis is systematic, and the codes will help you explore themes and patterns in the data. However, codes are not themes: they are just labels used to group similar types of data, developed to support your follow-up analysis.

There are two main approaches to coding. In *deductive coding*, the codes are decided upfront, before looking at the data, and may be based on your research problem phenomena, or may have emerged from your literature review, including codes possibly used in previous studies. In *inductive coding*, the codes emerge from the data and are not pre-defined. Deducting and inductive coding can also be combined by starting with a set of pre-defined codes then adding new codes as you review the data.

Whichever your approach, you should follow a multi-pass coding process. The first pass should consist of going through the whole data set in order to establish which codes to use. In the second pass, and any subsequent ones, you should apply the codes to the data bit by bit, say by line by line in a text, or frame by frame in a video, etc. In the second pass and subsequent passes, the initial codes are reviewed and may become more or less detailed.

There are various ways to choose codes. For instance, *in vivo* coding uses the exact language which occurs in the data: this is used, in particular, for participants' speech, especially when different languages are used. On the other hand, *descriptive*[•] coding uses words which encapsulate a general idea, such as 'sport' or 'running': this is particularly useful for non textual data, like images or videos.

Whichever codes you end up with, you should ensure they are properly defined, so that their are unambiguous and can be applied consistently. You should use a *codebook* for this purpose, which lists all the codes and their intended meaning, and that you can revisit and refine throughout the coding process.

The last step before detailed analysis is *code categorisation*, which is the process of reviewing what you have coded and organise it into categories. For instance, from codes such as 'football', 'tennis' and 'rugby' you may define a category 'sports'. In this way, you both organise your data and establish connections between codes and coded information.

- This is a very common approach, although there are others which you can research in the literature.

Both coding and categorisation are iterative processes which carry on until you reach saturation, that is no more is gained from further coding or categorisation. At this point, you can proceed with your chosen analysis method, whether content, thematic, narrative, discourse analysis or other, in order to identify patterns and themes, and provide your own interpretation of the data.

Coding and categorising are time consuming tasks, particularly if you have a large amount of text to code. In most research, coding data by hand is impractical and you should at least make use of a word processor, perhaps using colours and comments to code fragments of your text. Better still, you could make use of a bespoke qualitative data coding tool: many such tools are now available, some of which can also automate coding and categorisation to some extent.

Add URLs in the following question

Activity: Investigating tools for qualitative data coding

#33

Conduct a web search on tools which support qualitative data coding. List up to four which appear most commonly used. For each, indicate which coding features it offers and the extent it is freely available for students' research projects.

Discussion

Qualitative analysis tools are growing and changing rapidly, particularly due to the integration and exploitation of AI capabilities.

At the time of writing this book, the most used commercial products include NVivo, ATLAS.it and MAXQDDEThey all provide support for coding, with more or less extensive automation, alongside various other features such as data visualisation, statistical analysis, automatic transcripts generation from audio and video files, to name just a few. These commercial products are quite sophisticated with a steep learning curve and are usually quite expensive. They are also geared towards large research efforts, possibly by teams of researchers.

An increasing number of lighter, free products are also available. These include, for instance, Taguette, which supports manual coding and is both open source and free to use, or QDE Miner Lite, which is a free limited version of its full commercial release, and also supports manual coding. Such free products may be sufficient for Masters level research projects.

You may have found other similar tools.

0.4.3.2 Presenting qualitative data

While quantitative data can be summarised and presented using tables and charts, the same does not necessarily apply to qualitative data, which, due to their heterogeneous nature, cannot be easily set out in a standard manner.

Conveying the depth and richness of qualitative data in a succinct way is challenging, so that both selectivity and creativity are needed in presenting the data.

For textual data, like interview transcripts, verbatim quotations are often used to illustrate specific themes or points, or support certain conclusions. However, an excessive use of quotations will result in overlong accounts of the work, which may be difficult to follow or even obscure the main findings. Therefore it is important to select quotations which are particularly representative or poignant, avoiding verbose details that can be succinctly presented in the narrative around those quotations.

Diagrams, schematics or drawings can also be used effectively and imaginatively to present qualitative data and their analysis. Data visualisation is, in fact, a discipline in its own right^{*}, and some visualisation techniques can be applied to qualitative data.

Activity: Visualisation techniques for qualitative data

#34

Conduct a web search on techniques for visualising qualitative data. List the techniques you have found and what they are used for.

Discussion

You may have encountered some or all of the following techniques:

- diagrams and schematics, to convey complex processes or structures
- graphic timelines, to summarise key events and their order
- word clouds, to summarise emerging themes or concepts from text, and their relative frequencies
- mind maps, to visualise how different ideas relate or contribute to a central concept or topic
- heat maps, to highlight trends or differences in tabulated data
- icons, alongside brief descriptions, to represent and quickly identify specific concepts

- Edward Tufte is one of the most influential figures in this field. His books provide compelling examples on how to use visualisation to present and analyse highly complex data.

- bespoke drawings, for data which cannot be easily visualised using other standard techniques
- pie charts and bar charts, to summarise proportions and counts – which are actually quantitative, but may be the result of qualitative data analysis – of categorical data.

0.5 Writing up your analysis

In writing up your data analysis in your stage report or dissertation, you will need to decide:

- how to summarise your data and evidence. This will depend on their nature, and you will need to ensure that your summaries are appropriate to convey the essence of the evidence you have generated. In the previous sections, you have considered ways in which quantitative and qualitative data can be summarised using tables and visualisations. It may also be necessary for you to include sample raw data in an appendix.
- how to report findings. Your findings are your conclusions from your data analysis and should be reported as academic arguments which rely on the evidence you have generated.
- how to structure your narrative. Depending on your chosen research strategies and methods, different structures are possible. For instance, you may choose to start with a section which summarises all your evidence followed by one in which you analyse it, which may work well, for instance, for survey research. Alternatively, you could have separate sections each including a summary and analysis of a sub-set of your evidence: this may be appropriate for mixed methods research, with each section dealing with a different kind of data, or for design science research, with each section addressing a different design cycle. Whatever you choose, it is important that your report is effective in presenting your evidence and findings in a clear, rigorous and logical manner.

Activity: Writing up your analysis

#35

Consider the data you have collected and analysed so far. Note down how you are going to address each of the points above in your report. Write an outline of your analysis section.

Guidance

A good starting point is to consider how other researchers report their data analysis and findings. To this end, go back to some of the articles you have reviewed and consider their data analysis section and any related discussion. Ensure you select articles that apply similar collection and analysis methods to those in your research design, or deal with similar types of data.

0.6 Interpreting and evaluating data

Having generated and analysed a certain amount of data and evidence, it is time for you to start interpreting your findings in relation to your aim and objectives, and generally evaluate them in terms of their contribution to knowledge and possible limitations. This is a process you will repeat and complete in Stage 5, the concluding stage of your project, ending with your dissertation submission.

Interpreting your findings signifies addressing the following questions:

- What conclusions have you drawn from your data analysis?
- How do they relate to your aim and objectives?
- How do they relate to what you know from the literature?
- How do they relate to professional practice? (if applicable)
- Which new knowledge do they contribute?
- What do they fail to achieve?

Activity: Interpreting and evaluating your findings

#36

Consider your data analysis and based on it, address each of the above questions. Write down your responses, ensuring your arguments are well-formed, with explicit reference to evidence.

Guidance

Your interpretation and evaluation of findings will be, of course, limited by the data/evidence you have generated and analysed up to this point. You will revisit and expand this work in Stage 5 in order to complete your project.

0.7 Drafting an abstract for your project

An abstract is a common way to summarise academic research. Abstracts are an integral parts of all published academic articles – you will have encountered many abstracts while reviewing the literature. They are also very common in academic dissertations, therefore it is highly likely you will be required to include one at the beginning of yours.

An abstract provides a short summary of the whole research written for a specialist audience, that is you can assume that the reader has good knowledge of the topic and field of study. It should be a stand-alone item, so that it can be understood without reference to any other part of your dissertation.

Its content should convey succinctly the research problem, how and where it arises and its significance, the research aim and research design, key results obtained by the research, their evaluation and their implications for further research or professional practice.

Writing an abstract for your research is a good exercise, even if one is not needed for your dissertation, as it gives you an opportunity to write a logical argument that connects all key elements of your research. This can help you check that all the pieces fit together in a coherent manner. It is also something you can share with your supervisor and critical friends to communicate succinctly the essence of what you have done and achieved.

Activity: Drafting your abstract

#37

Write a draft abstract for your project, which should reflect your research progress to date.

Guidance

You should go back to some of the articles you have reviewed to consider the content and structure of their abstract. Choose a structure which may fit your project and write up your draft abstract accordingly. As your research is yet to be completed, you will not be able to write up the full abstract, but you should end up with a draft that you can easily complete by the end of your project.

0.8 Reflecting and reporting in Stage 4

It's time to write your Stage 4 report. As in the previous stages, before you do, it is worth reflecting on your work and learning in this stage.

Activity: Reflecting on your learning and practice

#38

As you did at the end of the previous stages, in this activity you are asked to stand back and reflect deeply on what you have learnt and done, the wider context of your work and your own attitude to it. Specifically, you are asked to think deeply about each of the following:

- your study this far
- the way you work
- the context of your research
- your feelings about your project

You should also think of any significant changes with respect to your reflection in the previous stages

Guidance

You should be accomplished at reflection by now. However, should you need to, you can refer back to the guidance to this activity in Stage 1, Section ??.

Your end-of-Stage 4 report will help you consolidate your work so far, adding yet another increment toward your full dissertation. We recommend you follow the guidance in Table 0.7 to write your report.

At the end of Stage 4, you should complete a report, extending that of Stage 3 and covering the work you have carried on in this stage. Its recommended structure and content are indicated in Table 0.7: much of the content should be carried forward from the previous stage.

Activity: Writing and assessing your report for Stage 4

#39

Using your word processor of choice, revise and expand your Stage 3 report by applying the structure and guidance in Table 0.7.

Assess your report by applying the criteria in Table 0.8. Revise and iterate until you are ready to move on.

Guidance

In completing your report, you should make good use of notes and summaries you wrote as part of the activities in this chapter. In evaluating your report, for each criteria, you should consider the related prompts, write down any further work needed for your next stage, and update your work plan and risk assessment table accordingly.

0.9 Takeaways

- Sampling is the process of selecting a sample from the population of interest, and is required in many research strategies. Many different approaches to sampling exist, depending on the nature and aim of your research.
- Questionnaires are common tools for data generation. Good questionnaire design relies on a wide range of considerations (see Section ??).

Table 0.7: Report structure and guidance

Report template	Guidance
Proposed title	Your title should continue to capture succinctly your research problem and aim. <i>It is likely this is the same as, or very similar to, that in Stage 3</i>
Abstract	You should include your draft abstract providing a succinct account of your research to date
Sect 1 - Introduction 1.1 Background to the research 1.2 Justification for the research 1.3 fitness of the research	This section should continue to provide an introduction to your research topic in its wider context (as background) and your justification of why the research is worth pursuing. Its purpose is to introduce and justify your intended research in overview, before entering the detailed work of the subsequent sections. It should be well argued and supported by appropriate citations. In this section, you should also argue how the research fits within the scope of your qualification, and meets any other personal, professional or organisational criteria. <i>You may review this section from Stage 1 to reflect your growing understanding of the topic in context derived from your literature review.</i>
Sect 2 - Literature review 2.1 Review of existing relevant knowledge 2.2 Critical summary, including knowledge gap to be addressed by the research	Your review should provide a critical account of your in-depth engagement with the academic (and other) relevant literature, including identifying key trends, ideas and possible knowledge gaps. Most of your citations should point to academic articles. Your critical summary should highlight key insights from your review and provide a strong justification for your proposed research. Both coverage and depth of your review matter. You should ensure that your review is well structured, with a logical narrative flow and your arguments are well supported by evidence
Sect 3 - Research definition 3.1 Problem statement 3.2 Aim, objectives, tasks and deliverables 3.3 Knowledge contribution	You should ensure that your research problem is well articulated and appropriate for your course and your personal and professional circumstances, that your aim and objectives are consistent with research problem, that tasks and deliverables break down your objectives appropriately and are clearly related to your chosen research methods, and that the intended knowledge contribution of your research is clearly articulated
Sect 4 - Research design 4.1 Evidence and data 4.2 Research strategy and methods 4.3 Research procedures 4.4 Ethical, legal and EDI considerations	This section should demonstrate your critical engagement with all elements of research design, including a detailed account of the data and evidence needed in your research, the research methods and research strategies chosen, with justification, and applied within your project. Your account should be supported by a clear rationale and insights from the related literature, and appropriately justified in relation to your research problem, aim and objectives. It should also demonstrate your careful consideration of ethical and legal matters, and that your research complies with your

Table 0.8: Criteria for reviewing your research proposal

Criteria	Prompts
Completeness	Are all sections included and their content complete? What is missing?
Academic writing	Have you applied good academic writing practices throughout? Which main issues do you still have to address?
Logical structure and flow	Have you structured your writing appropriately to ensure a logical flow of arguments? Which restructuring may be needed?
Supporting evidence	Are your key arguments supported by appropriate references or other evidence? Which further evidence is needed?
Citation and reference style	Do all your citations and references comply with the required bibliographical style?
Avoiding plagiarism	Have you acknowledged the work of others and distinguished it from your own appropriately?
Grammar and spelling	Have you proof-read your report carefully to remove all typos and grammatical errors?

- When a large amount of raw data is collected, it is important to devise appropriate ways to store and organise them, paying particular attention to backing them up and protecting personal data.
- Tables are common ways to organise and present data, and a good starting point for data analysis. Pivot, frequency and contingency tables are commonly used in research.
- Descriptive statistics is used to describe data, with various attributes of data sets defined and calculated, such as centrality, dispersion or skewness. Charts are often used to visualise such attributes.
- Inferential statistics is used to make predictions from data, specifically to establish whether patterns or effects observed on sample data can be inferred for the whole population from which the sample was taken.
- Statistical tests are used to establish the statistical significance of observations on a sample in relation to the whole population. They are used both for comparing data to set values and to establish relationships between variables. Many statistical tests exist.
- Coding is the first step in qualitative analysis, and is the process of assigning labels to extracts from a qualitative data set to allow a systematic follow-up analysis. Different approaches to coding exist.

- Qualitative data are heterogeneous in nature, so that they cannot be easily set out in a standard manner. Many different, often bespoke, approaches to present and visualise qualitative data have been proposed in the literature.
- In writing up your data analysis you must decide how to summarise your data, how to report your findings and how to structure your narrative.
- Interpreting your findings means to indicate what you can conclude from the data, how that relate to your aim and objectives, and which new knowledge it contributes.
- An abstract is a short summary of your whole dissertation, written for a specialist audience as a stand-alone piece, that is understandable without reference to any other part of your dissertation.
- The template provided can help you structure your Stage 4 report.